

chapter IV



service architecture



data partitioning



multitenancy



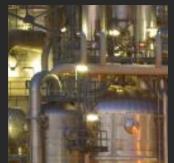
DevOps

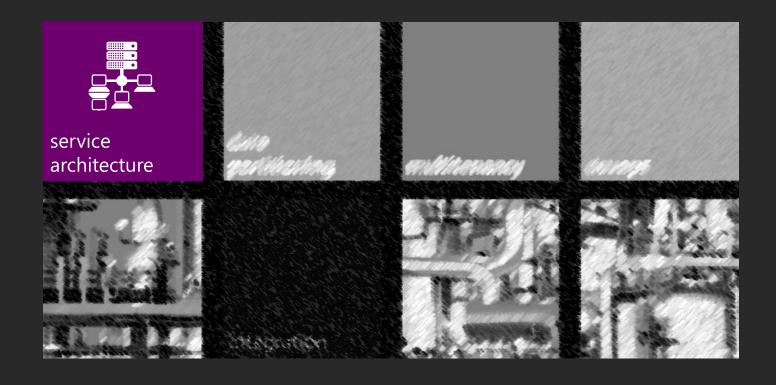




integration







service architecture

Device Client/Browser

REST/WS-*/...
Service Interface

Process Layer

Business Layer

Data Layer

service architecture

Device Client/Browser

REST/WS-*/... Service Interface **Process Layer Business Layer** Data Layer

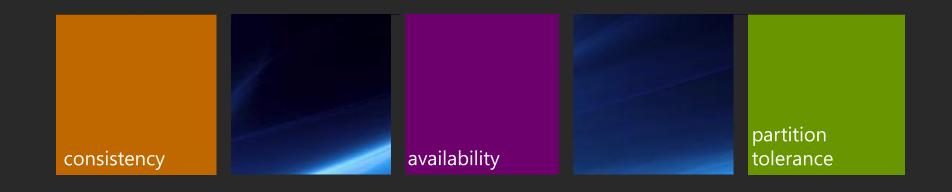
REST/WS-*/... Service Interface **Process Services Business Services Data Services**

each instance contains all layers inproc communication

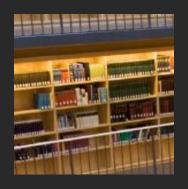
instance per layer, communication through queues



cap theorem



data partitioning



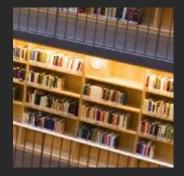
traditional reasons

data volume (too many bytes)

work load (too many transactions/second)

new 'cloud focused' reasons cost (using different cost storage)

elasticity (just in time partitioning for high load periods)



horizontal partitioning

First Name	Last Name	Email	Thumbnail	Photo
David	Alexander	davida@contoso.com	3kb	3MB
Jared	Carlson	jaredc@contoso.com	3kb	3MB
Sue	Charles	suec@contoso.com	3kb	3MB
Simon	Mitchel	simonm@contoso.com	3kb	3МВ
Richard	Zeng	richardz@contoso.com	3kb	3MB

horizontal partitioning (sharding)



spread data across similar nodes

achieve massive scale out (data and load)





intra-partition queries are simple

cross-partition queries are harder



vertical partitioning

First Name	Last Name	Email	Thumbnail	Photo
David	Alexander	davida@contoso.com	3kb	3МВ
Jared	Carlson	jaredc@contoso.com	3kb	3MB
Sue	Charles	suec@contoso.com	3kb	3МВ
Simon	Mitchel	simonm@contoso.com	3kb	3МВ
Richard	Zeng	richardz@contoso.com	3kb	3МВ
SQL Azure			Tables	BLOBS

vertical partitioning



place frequently queried data in more 'expensive' indexed storage



retrieving
a whole row requires
>1 query

spread data across dis-similar nodes



place large data in 'cheap' binary storage

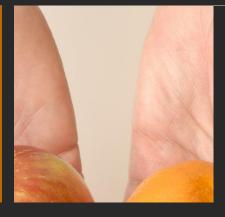


hybrid partitioning

First Name	Last Name	Email	Thumbnail	Photo
David	Alexander	davida@contoso.com	3kb	3МВ
Jared	Carlson	jaredc@contoso.com	3kb	3МВ
Sue	Charles	suec@contoso.com	3kb	3МВ
Simon	Mitchel	simonm@contoso.com	3kb	3МВ
Richard	Zeng	richardz@contoso.com	3kb	3МВ

tables != rdbms

cross partition queries are resource intensive



aggressive data duplication can save money and boost performance

storage is cheap



goal: To be able to include Partition Key in all queries

Tweet

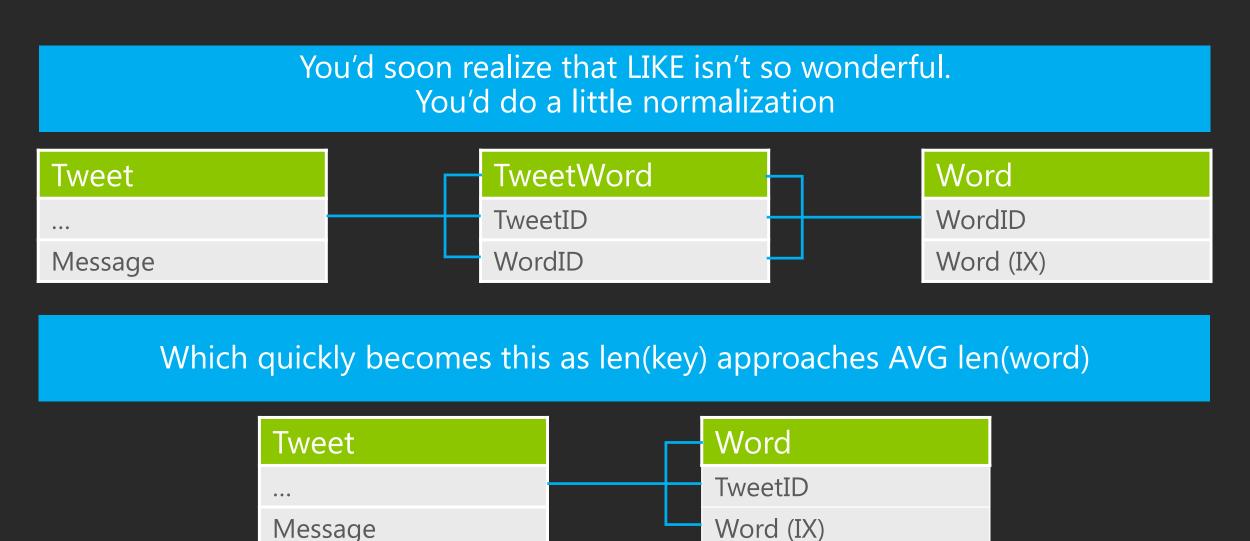
TweetID

UserID

DateTimeStamp

Message

With an RDBMS you'd probably start something like this:
SELECT * FROM Tweet WHERE
Message Like %SearchTerm%



With Tables we go the whole way



TweetID (RK)

UserID (PK)

DateTimeStamp

Message

Worker Role Creates

TweetIndex

TweetID (RK)

UserID

DateTimeStamp

Message

Word (PK)

GET All Entities in Partition 'DavidA' from Tweet
GET All Entities in Partition 'Foo' from TweetIndex

We may create multiple indexes



TweetID (RK)

UserID (PK)

DateTimeStamp

Message

Worker Role Creates

MentionIndex

TweetID (RK)

UserID

DateTimeStamp

Message

MentionedUserID (PK)

GET All Entities in Partition 'DavidA' from MentionIndex

traffic management fundamentals



performance

directs the user to the "best"/"closest" deployment



direct the user to the "best" deployment between US South and West Europe



failover

one deployment is primary

traffic is redirected to another deployment if the primary goes down

example:

all traffic is directed to US North; if it goes down, send all traffic to US South



geomapping

allows users from defined geographic locations to be directed to particular deployment

example:

all users from US -> US North, all users from Asia -> US North, all users from Europe -> West Europe



ratio

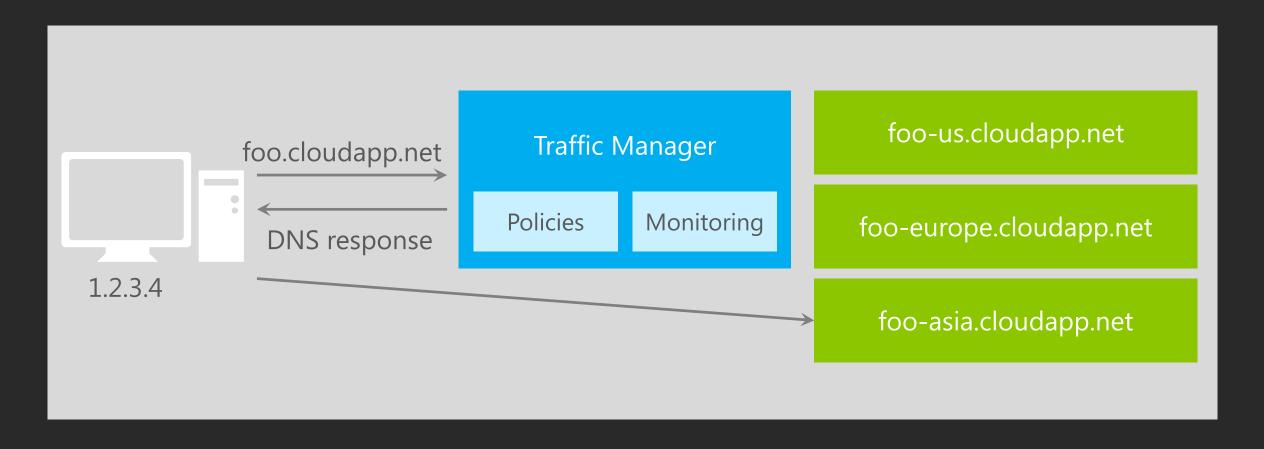
sends traffic to different deployments based on fixed ratio (N/M)

example:

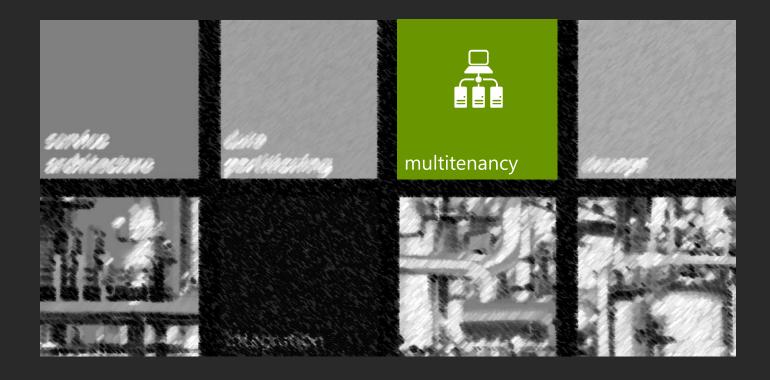
direct 20% of user traffic to US South and 80% to US North.

Windows Azure Traffic Manager

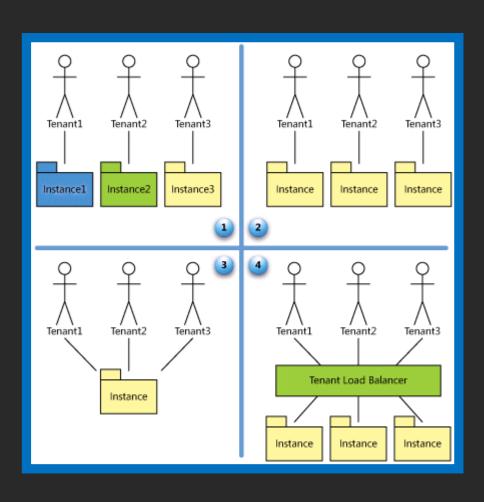
Direct users to the service in the closest region with the Windows Azure Traffic Manager







multitenancy





continues deployment

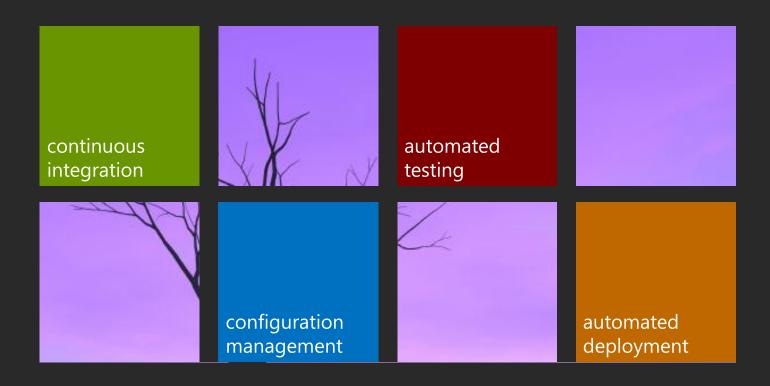


http://code.flickr.com/

devops



devops



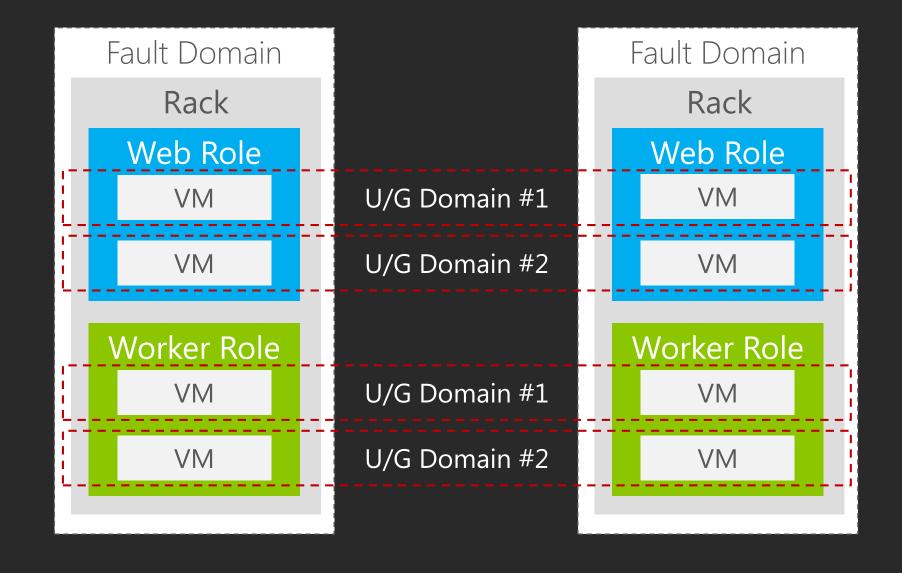
deployment/release strategies



upgrades in Windows Azure



fault and upgrade domains

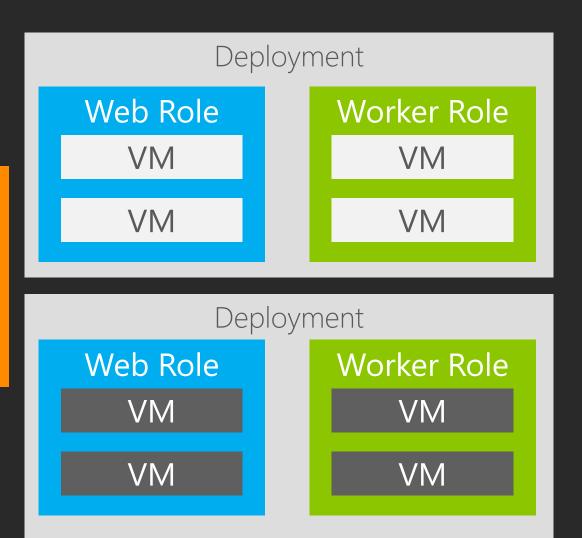


VIP swap

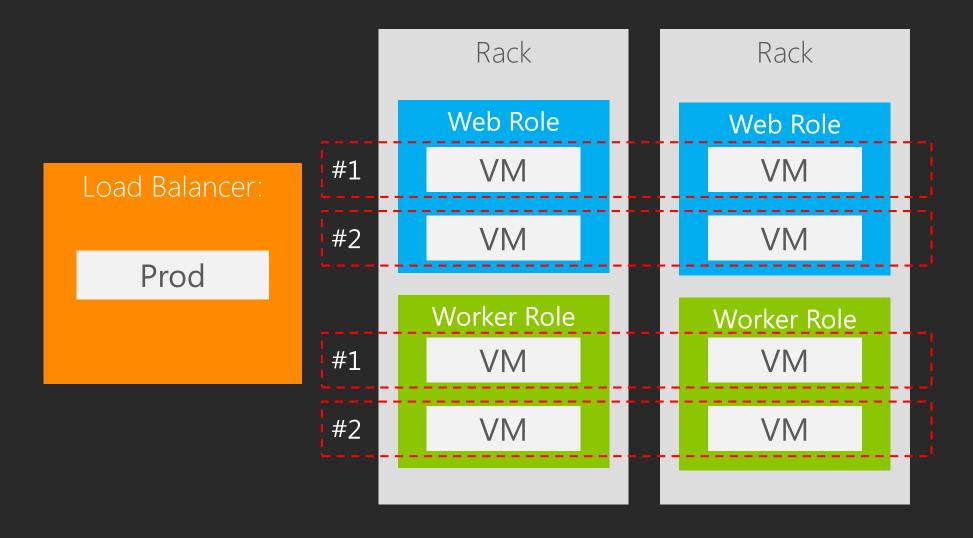
Load Balancer:

Prod

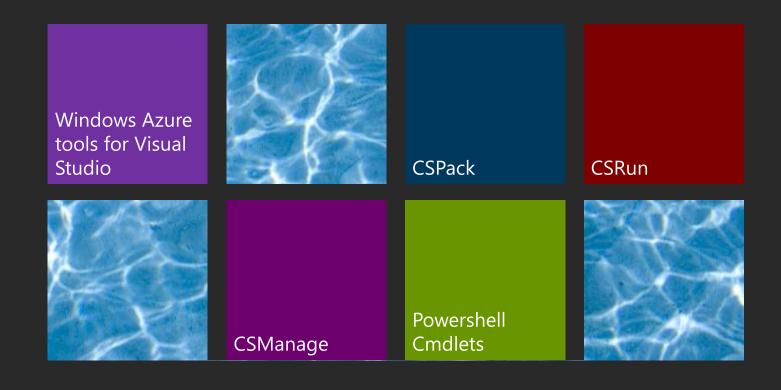
Stage



in place upgrade



Windows Azure deployment tools

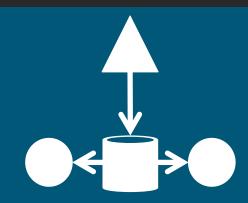




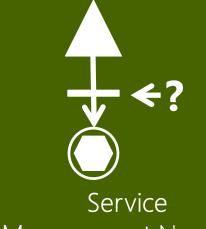
Windows Azure Service Bus



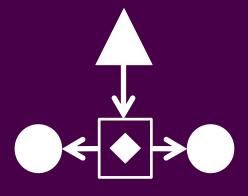
Connectivity
Service Relay
Protocol Tunnel
Eventing, Push



Messaging
Queuing
Pub/Sub
Reliable Transfer



Service
Management Naming,
Discovery
Monitoring



Integration
Routing
Coordination
Transformation

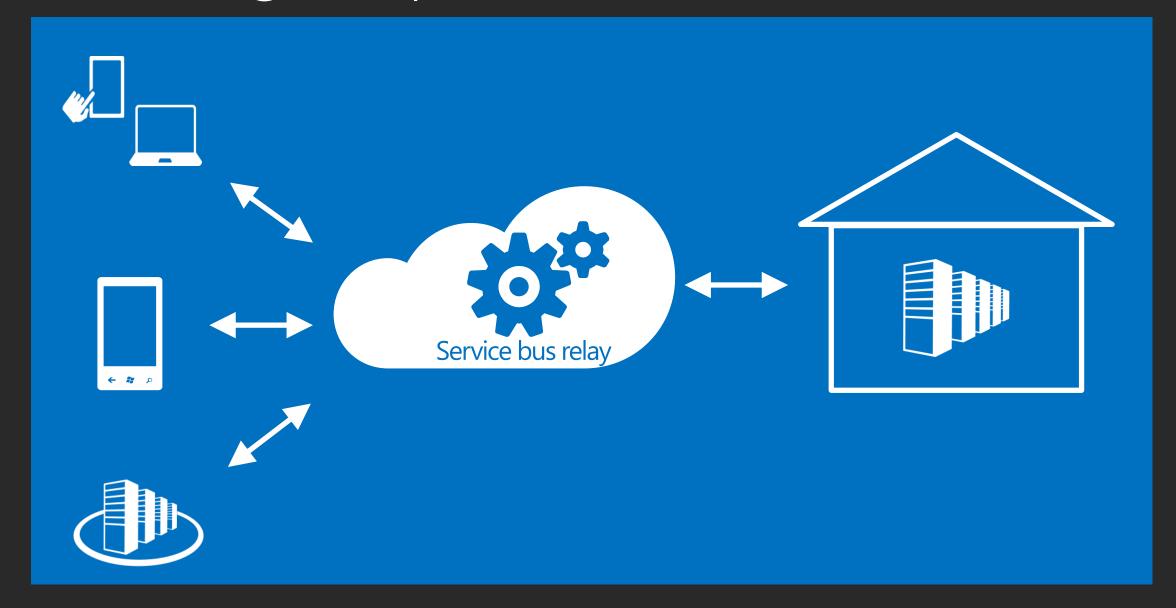
Rich options for interconnecting apps across network boundaries

Reliable, transactionaware cloud messaging infrastructure for business apps

Consistent management surface and service observation capabilities

Content-based routing, document transformation, and process coordination

accessing on-premise services





conclusion



cloud style computing designs for scale out

scale out requires partitioning





having redundant/duplicated data is ok

continuous delivery is a key asset



chapter V

closing

