



cloud
strategy
day


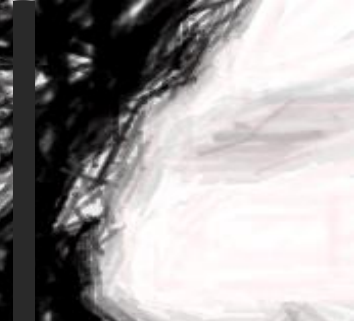


“How the Microsoft cloud enables the best App experiences across devices”

Beat Schwegler – beatsch@microsoft.com
Director, Platform Strategy Group, Microsoft Corp.

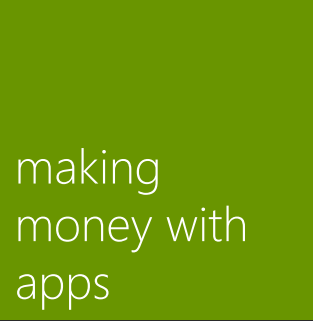
agenda




Microsoft's
devices and
services
strategy



making
money with
apps



cloud enabled
app scenarios
with Windows
Azure



key
architectural
considerations

chapter IV



service
architecture



data
partitioning



multitenancy



DevOps



integration





service
architecture

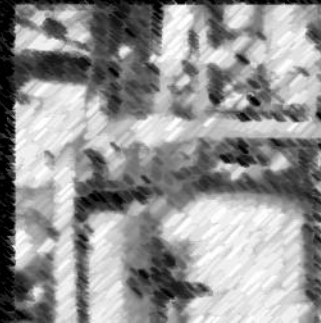
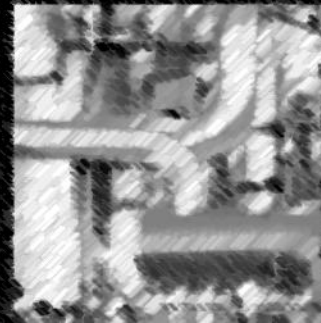
data
partitioning

multitenancy

design



integration



service architecture

Device Client/Browser

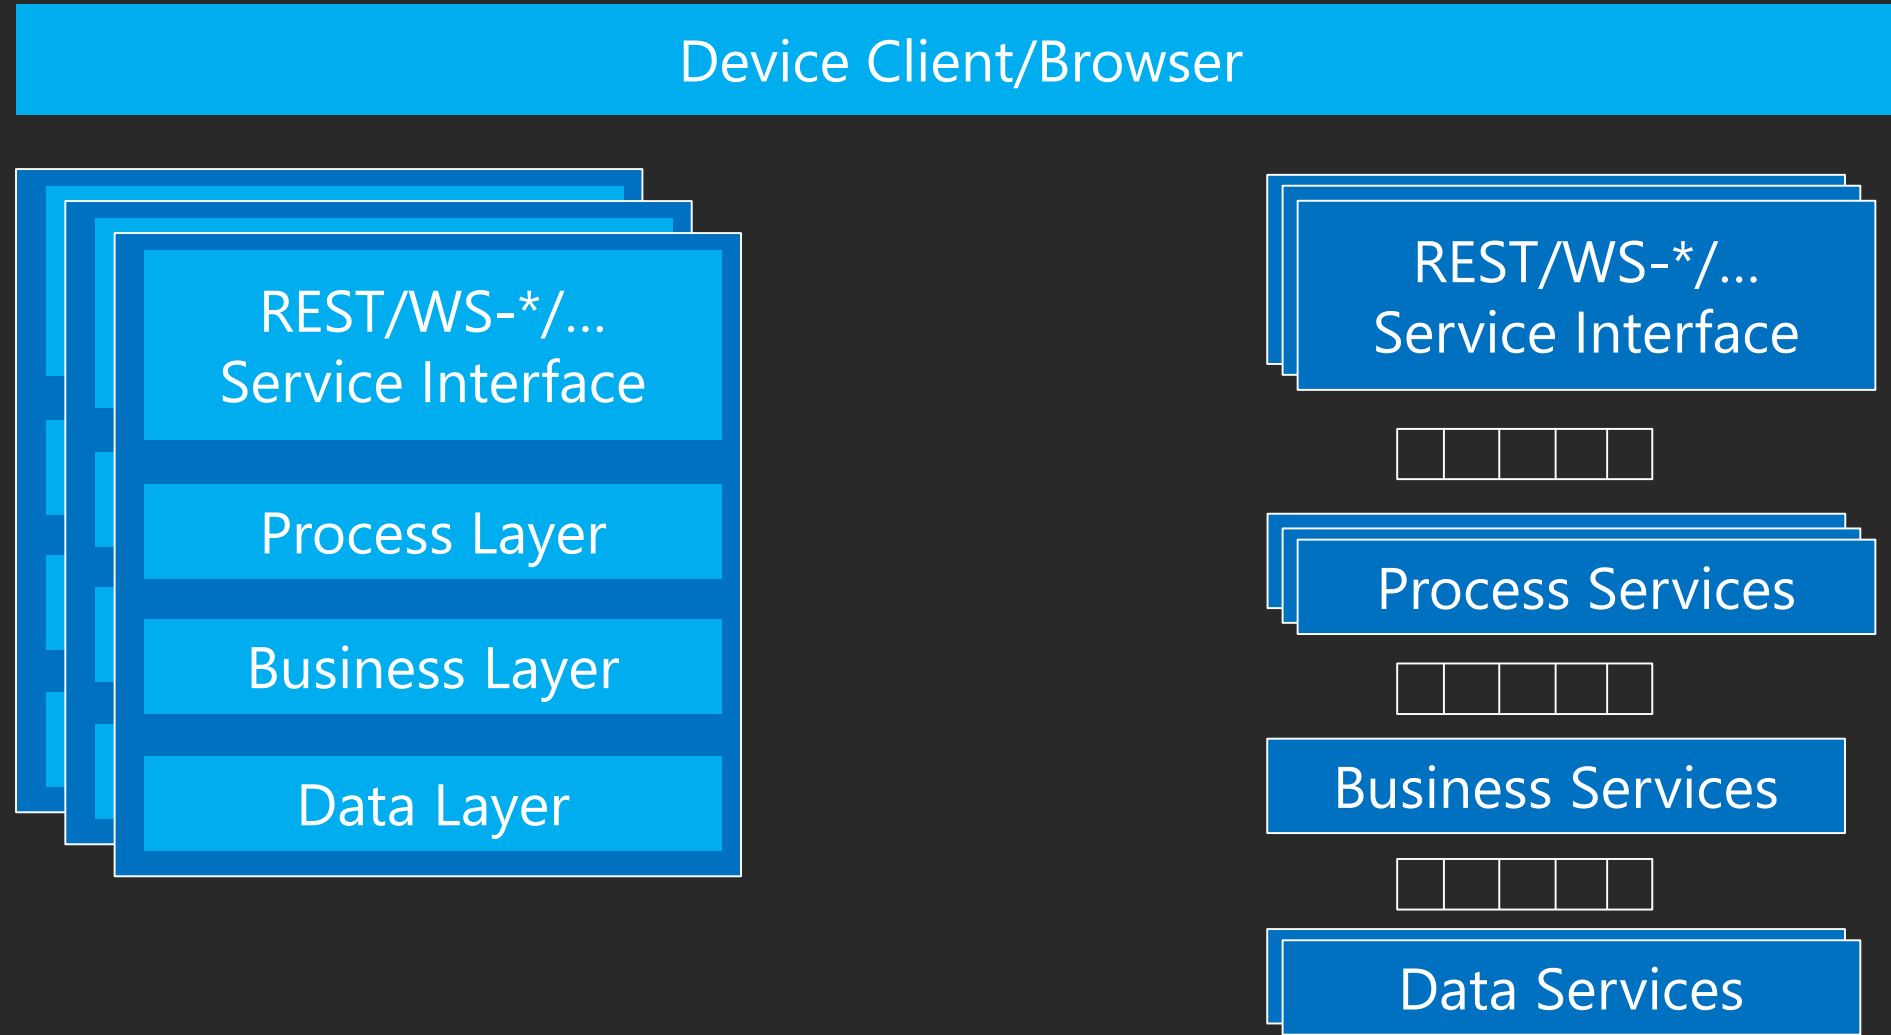
REST/WS-*/...
Service Interface

Process Layer

Business Layer

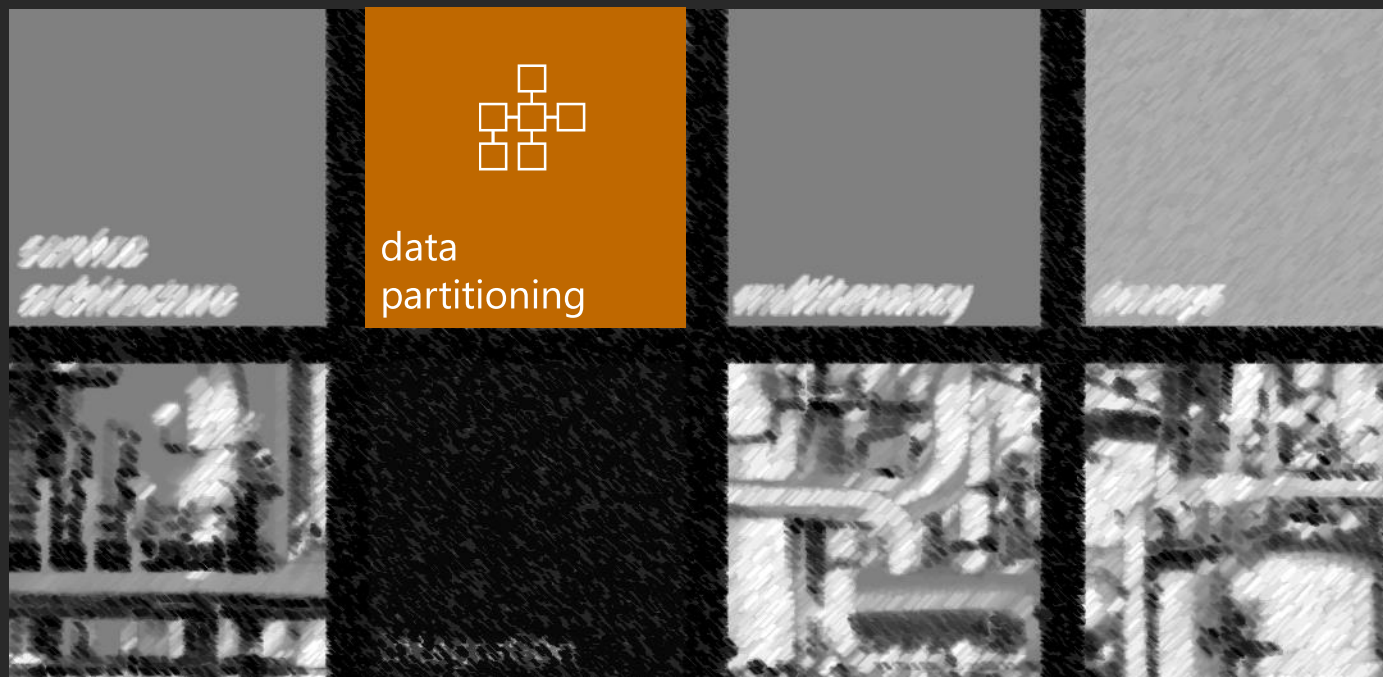
Data Layer

service architecture

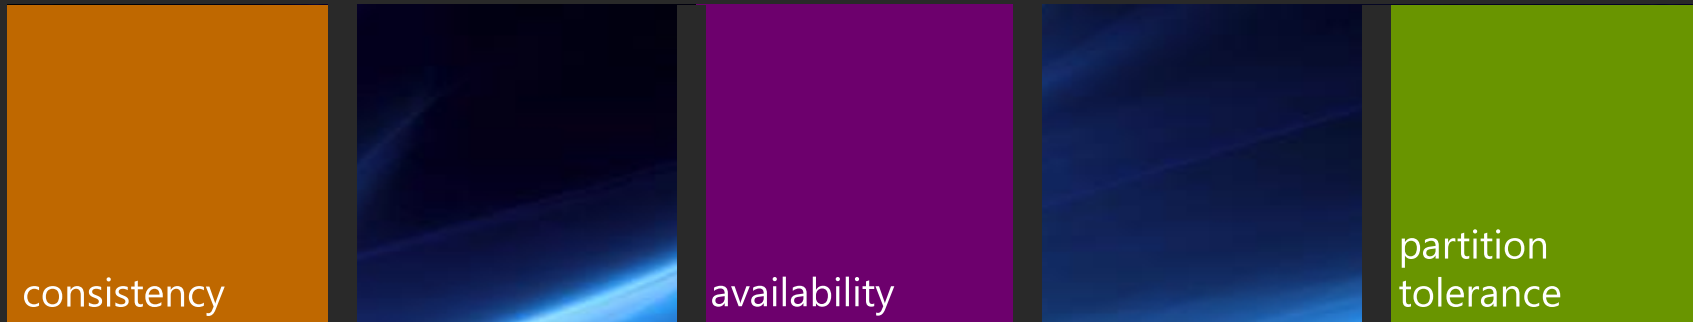


each instance contains all layers
inproc communication

instance per layer, communication through queues



cap theorem



data partitioning



traditional
reasons

data volume (too many bytes)

work load (too many
transactions/second)

new 'cloud
focused' reasons


cost (using different cost storage)

elasticity (just in time partitioning
for high load periods)



horizontal partitioning

First Name	Last Name	Email	Thumbnail	Photo
David	Alexander	davida@contoso.com	3kb	3MB
Jared	Carlson	jaredc@contoso.com	3kb	3MB
Sue	Charles	suec@contoso.com	3kb	3MB
Simon	Mitchel	simonm@contoso.com	3kb	3MB
Richard	Zeng	richardz@contoso.com	3kb	3MB



The diagram illustrates horizontal partitioning of a table. The table is divided into four partitions based on the 'First Name' column. Each partition is stored on a separate server. The partitions are color-coded: blue for David, orange for Jared and Sue, purple for Simon, and red for Richard. Arrows point from the 'First Name' column of each row to a server icon below. The servers are represented by a blue server rack and a blue database cylinder.

horizontal partitioning (sharding)



spread data across similar nodes

achieve massive scale out (data and load)



intra-partition queries are simple

cross-partition queries are harder



vertical partitioning

First Name	Last Name	Email	Thumbnail	Photo
David	Alexander	davida@contoso.com	3kb	3MB
Jared	Carlson	jaredc@contoso.com	3kb	3MB
Sue	Charles	suec@contoso.com	3kb	3MB
Simon	Mitchel	simonm@contoso.com	3kb	3MB
Richard	Zeng	richardz@contoso.com	3kb	3MB



SQL Azure



Tables



BLOBS

vertical partitioning



place frequently
queried data in more
'expensive' indexed
storage



retrieving
a whole row requires
>1 query

spread data across
dis-similar nodes



place large data in
'cheap' binary
storage



hybrid partitioning

First Name	Last Name	Email	Thumbnail	Photo
David	Alexander	davida@contoso.com	3kb	3MB
Jared	Carlson	jaredc@contoso.com	3kb	3MB
Sue	Charles	suec@contoso.com	3kb	3MB
Simon	Mitchel	simonm@contoso.com	3kb	3MB
Richard	Zeng	richardz@contoso.com	3kb	3MB

Diagram illustrating hybrid partitioning. The table is partitioned by Last Name. The 'First Name', 'Last Name', and 'Email' columns are grouped into two partitions: one for 'Alexander', 'Carlson', and 'Charles' (blue outline), and another for 'Mitchel' and 'Zeng' (orange outline). The 'Thumbnail' column is grouped into one partition (purple outline), and the 'Photo' column is grouped into another partition (red outline). Arrows indicate that the 'First Name', 'Last Name', and 'Email' partitions are stored in two separate database instances (server and disk icons), while the 'Thumbnail' and 'Photo' partitions are stored in two separate database instances.

tables != rdbms

cross partition
queries are
resource
intensive



aggressive data
duplication can
save money
and boost
performance

storage is
cheap



goal: To be able
to include
Partition Key in
all queries

e.g. tweet storage

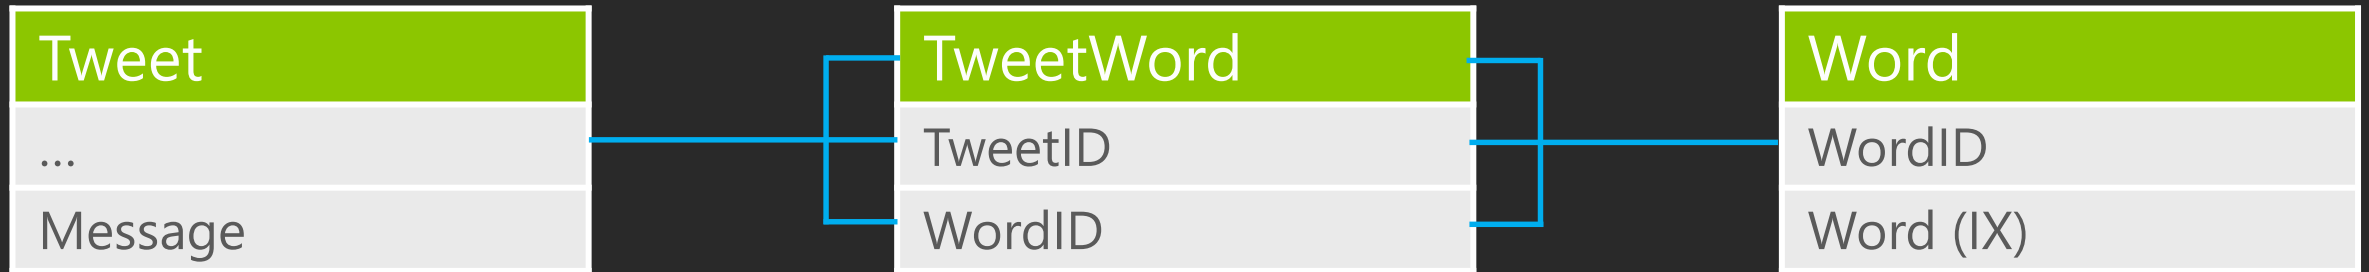
Tweet
TweetID
UserID
DateTimeStamp
Message

With an RDBMS you'd probably start something like this:

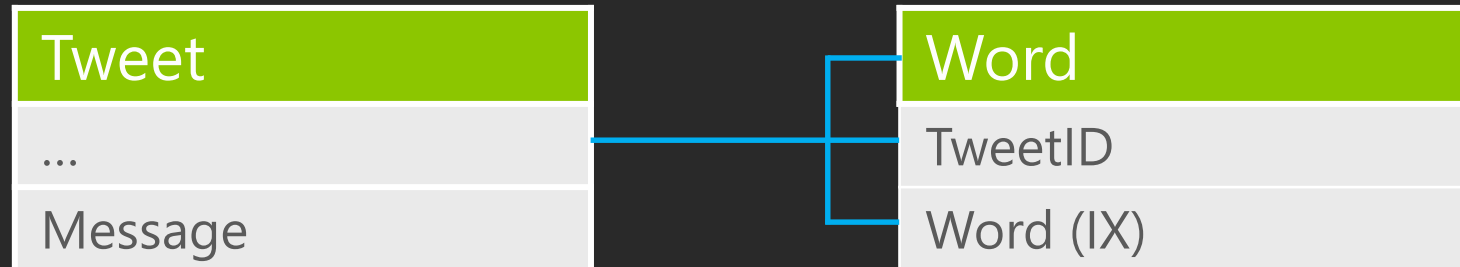
```
SELECT * FROM Tweet WHERE  
Message Like %SearchTerm%
```

e.g. tweet storage

You'd soon realize that LIKE isn't so wonderful.
You'd do a little normalization



Which quickly becomes this as $\text{len}(\text{key})$ approaches $\text{AVG len}(\text{word})$



e.g. tweet storage

With Tables we go the whole way

Tweet
TweetID (RK)
UserID (PK)
DateTimeStamp
Message

Worker Role Creates

TweetIndex
TweetID (RK)
UserID
DateTimeStamp
Message
Word (PK)

GET All Entities in Partition 'DavidA' from Tweet
GET All Entities in Partition 'Foo' from TweetIndex

e.g. tweet storage

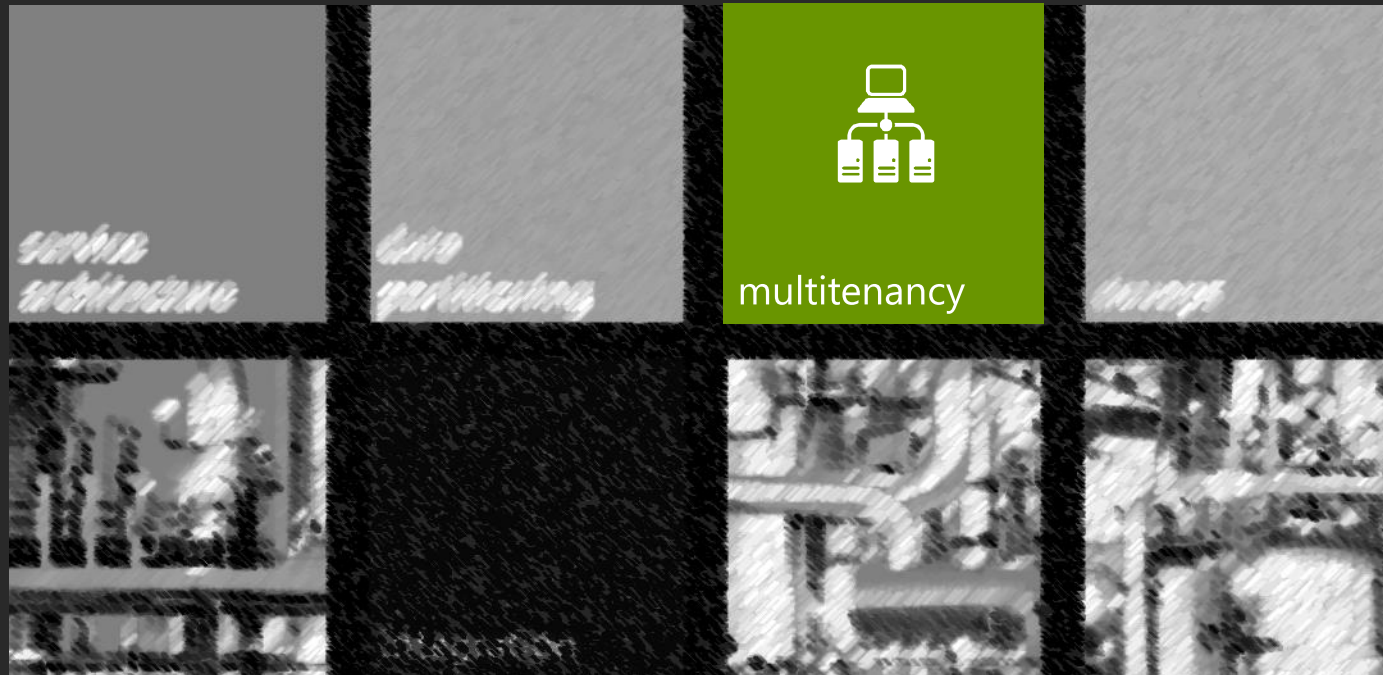
We may create multiple indexes

Tweet
TweetID (RK)
UserID (PK)
DateTimeStamp
Message

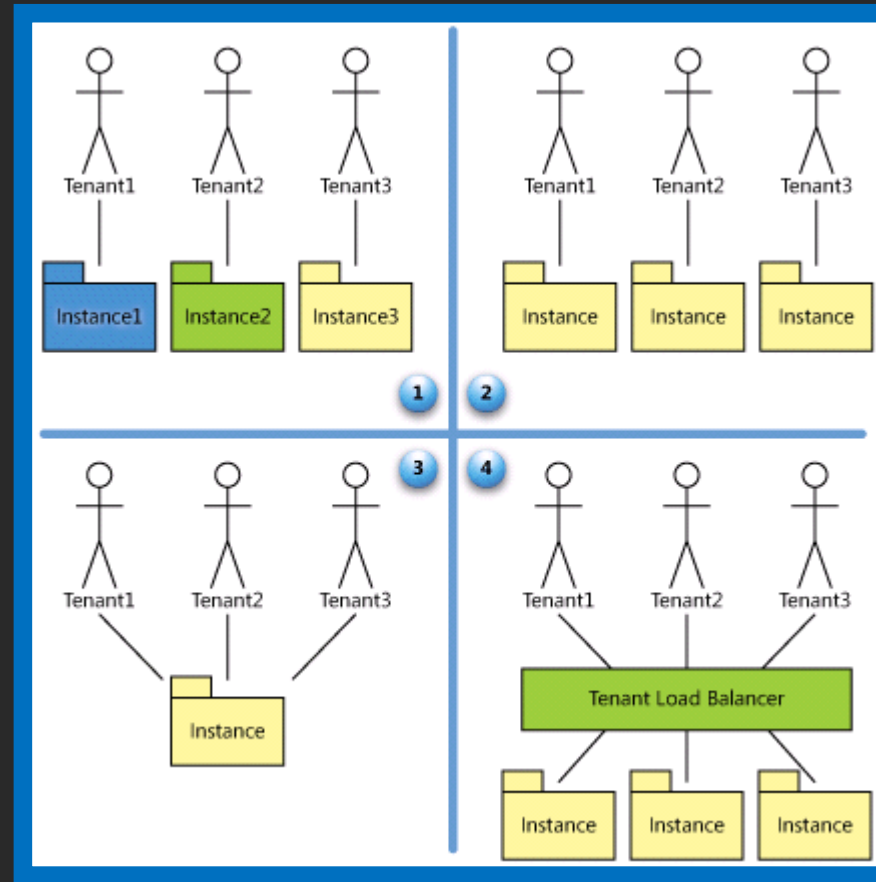
Worker Role Creates

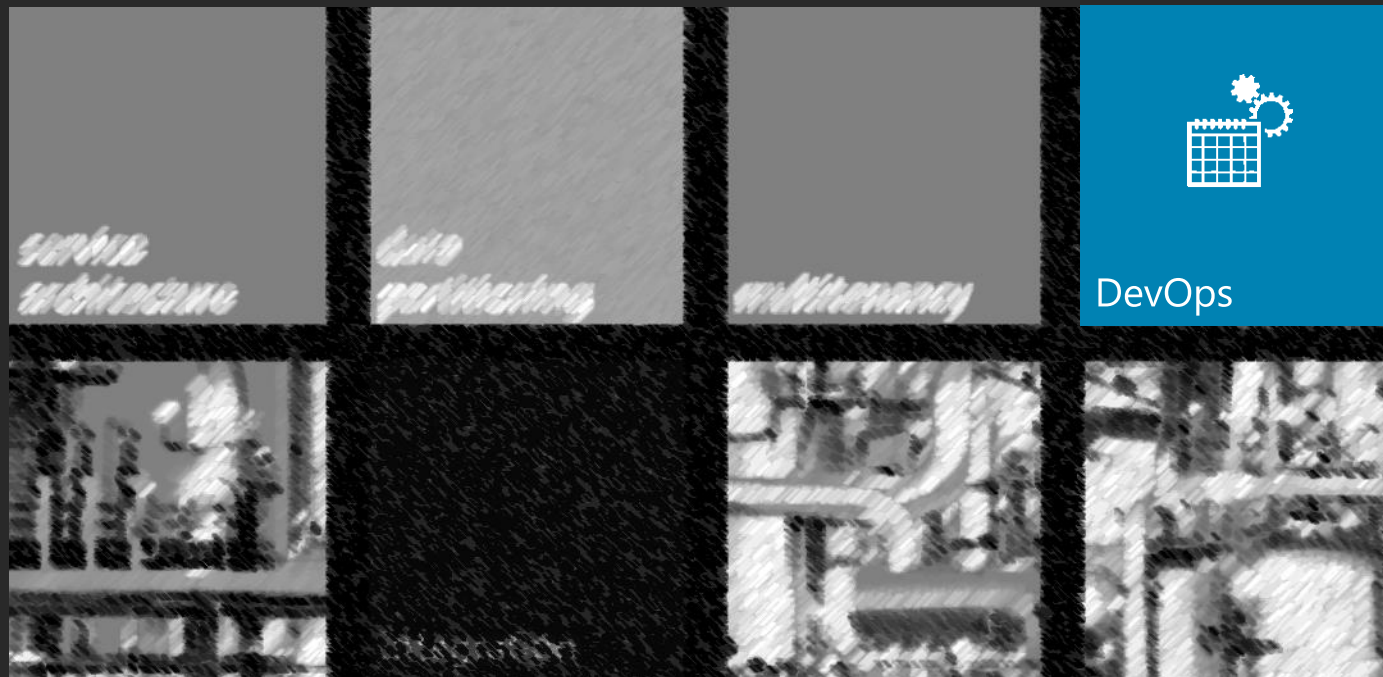
MentionIndex
TweetID (RK)
UserID
DateTimeStamp
Message
MentionedUserID (PK)

GET All Entities in Partition 'DavidA' from MentionIndex




multitenancy





continues deployment

FEATURING



Flickr was last deployed 2 days ago, including 4 changes by 1 person.

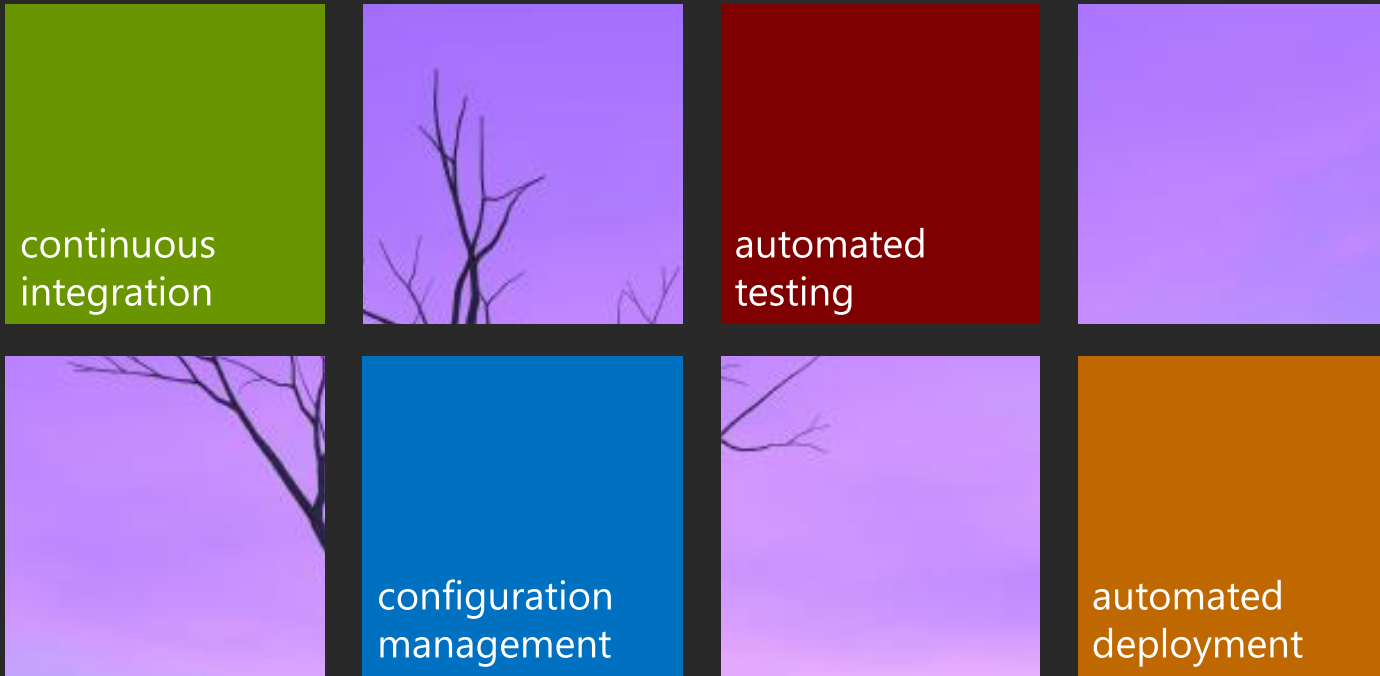
In the last week there were 77 deploys of 440 changes by 18 people.

<http://code.flickr.com/>

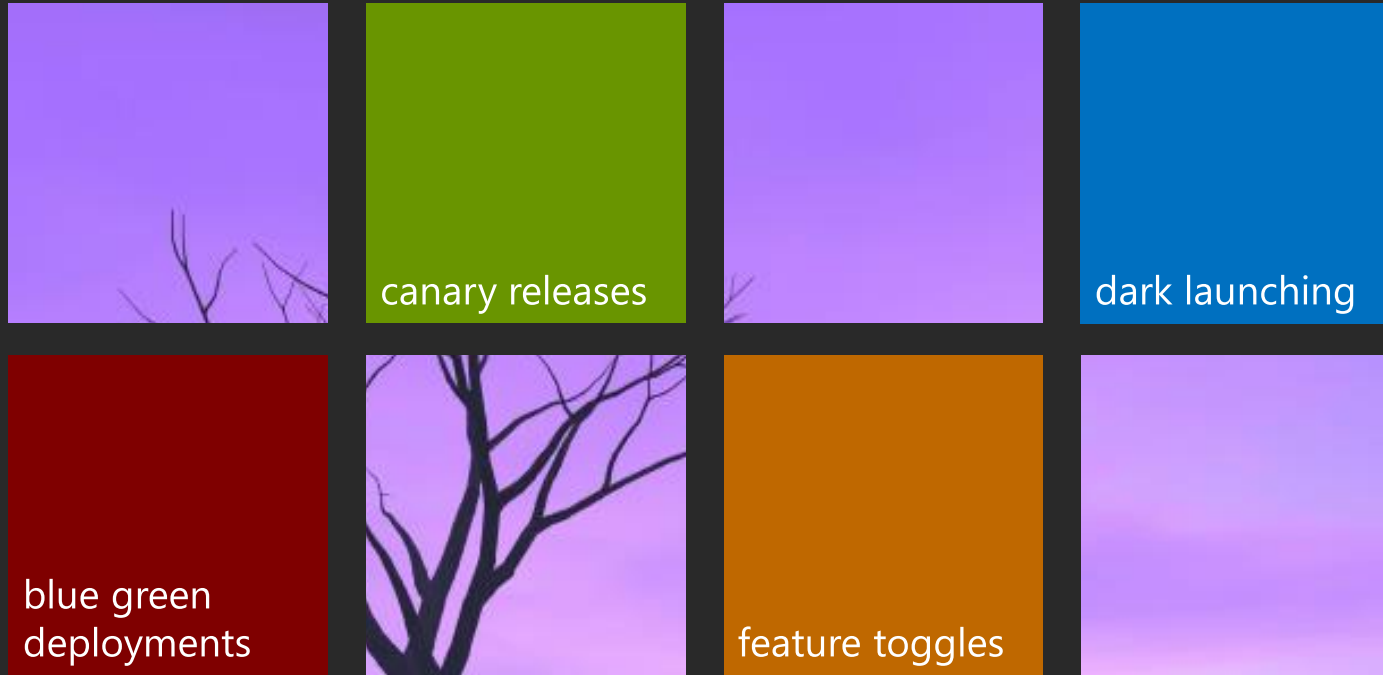
devops



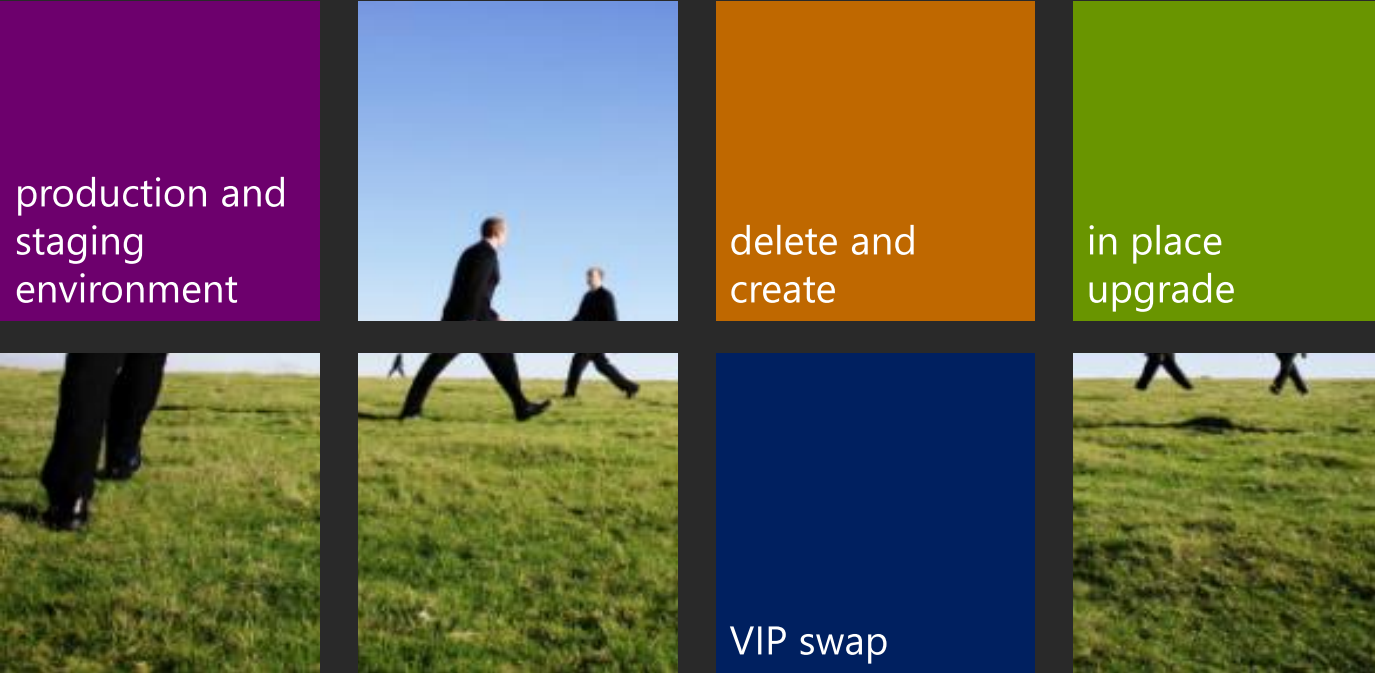
devops



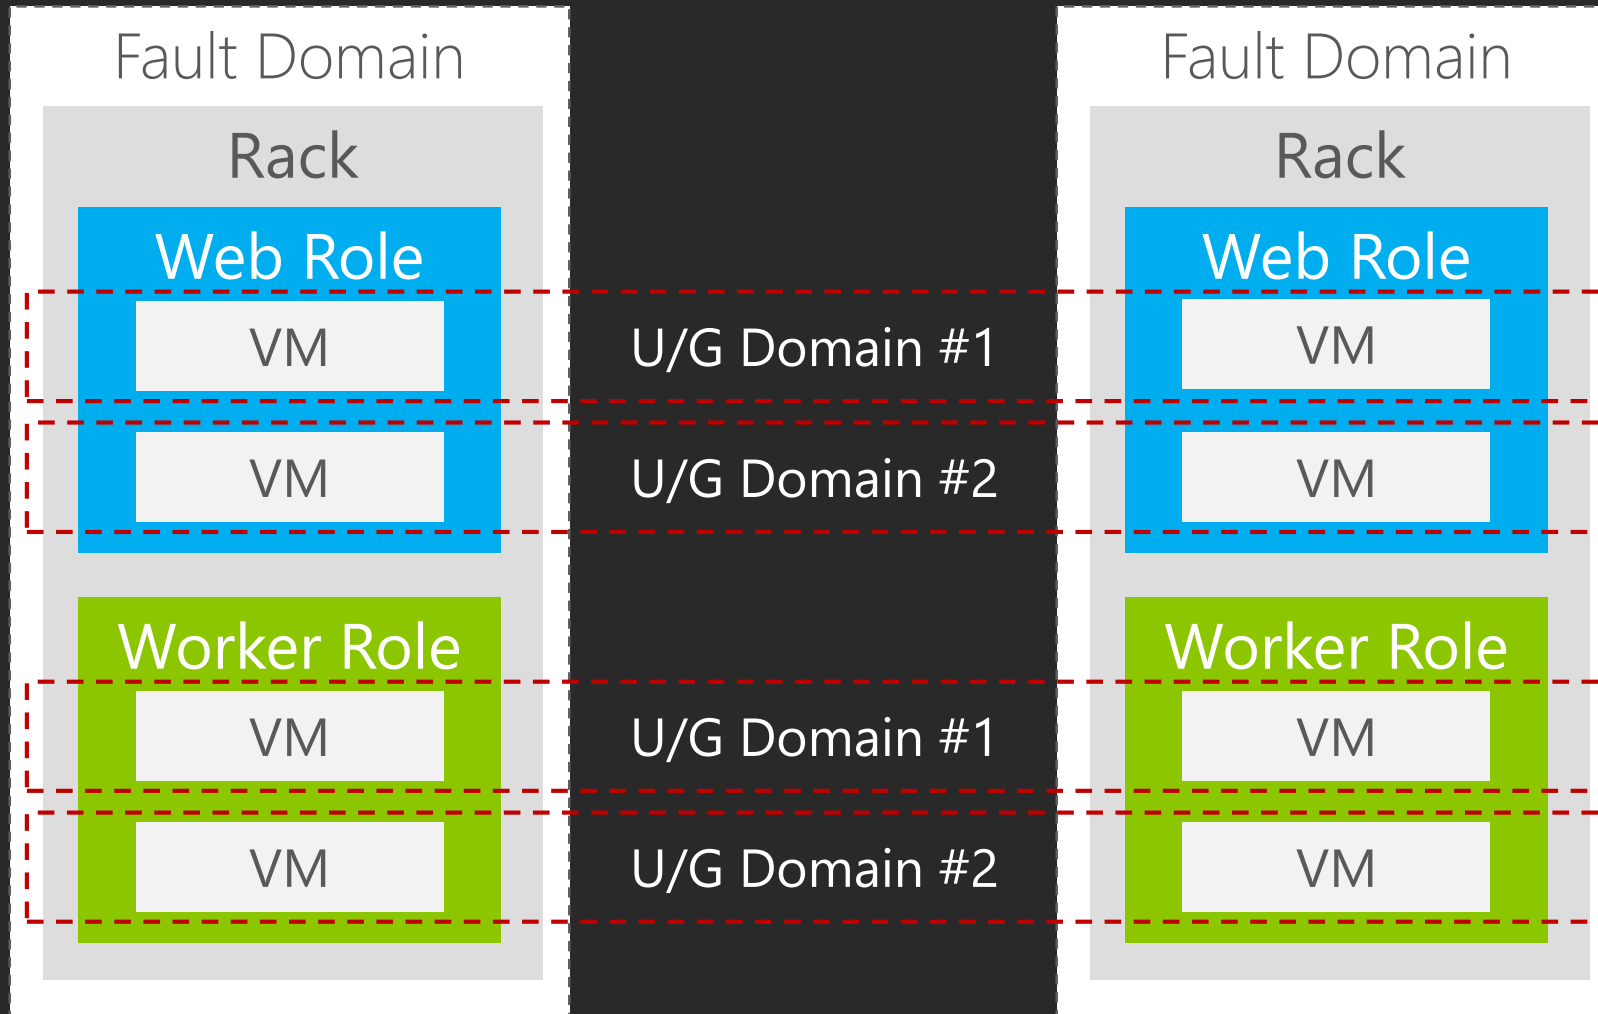
deployment/release strategies



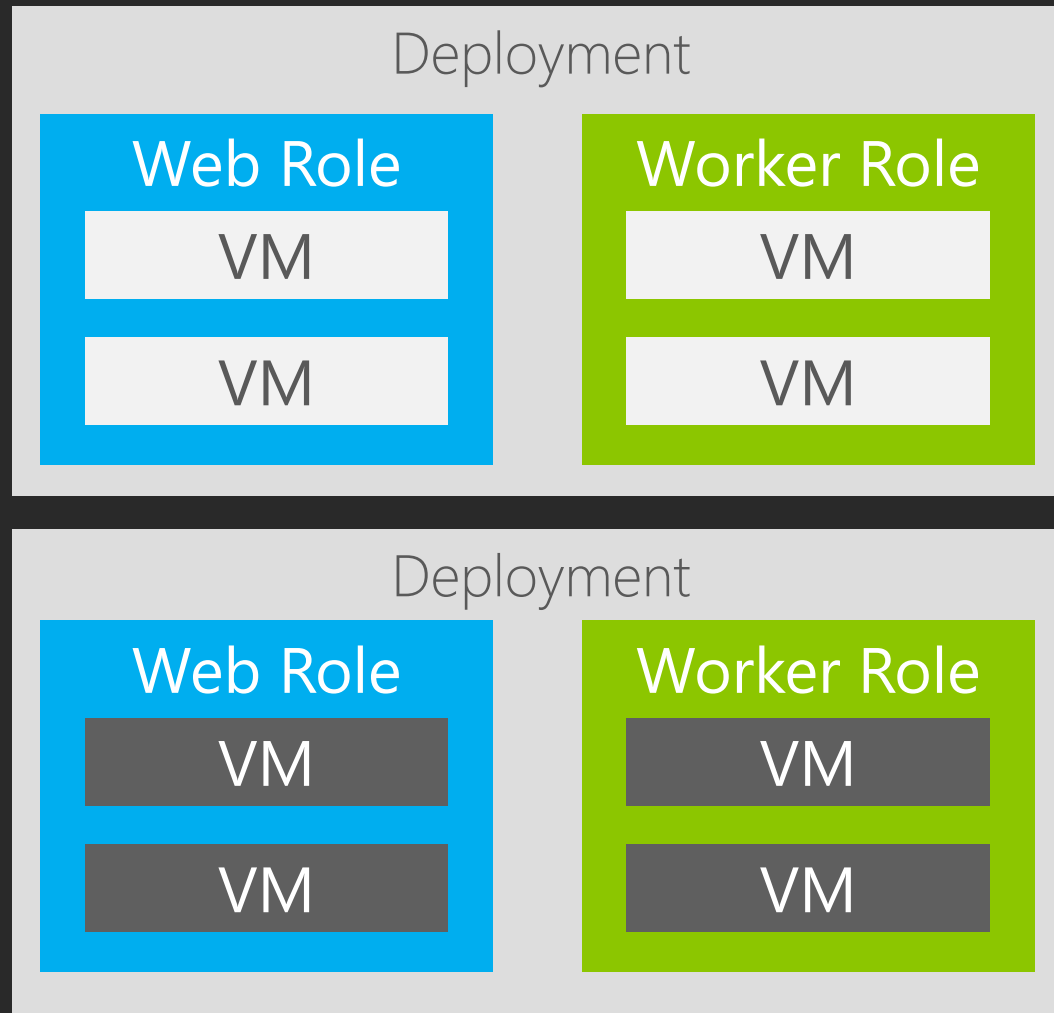
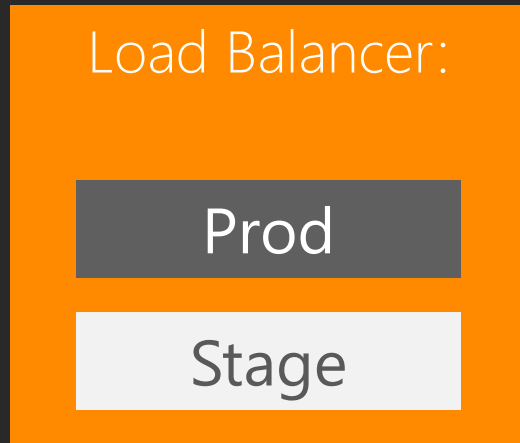
upgrades in Windows Azure



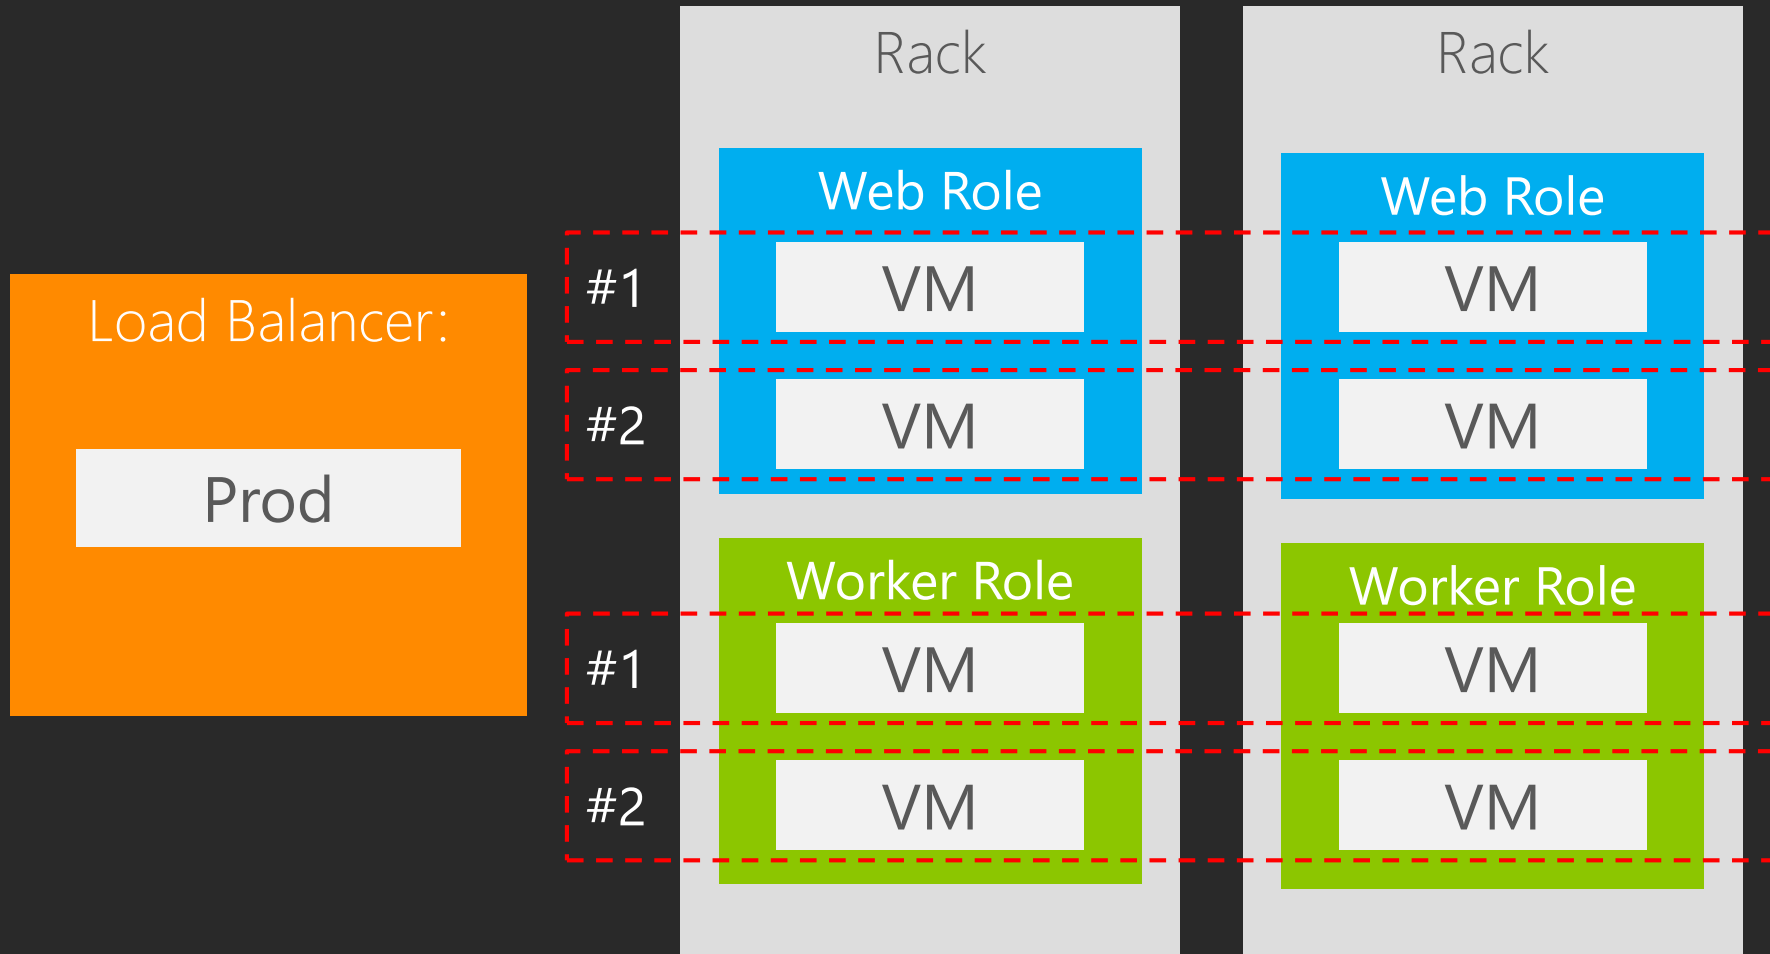
fault and upgrade domains



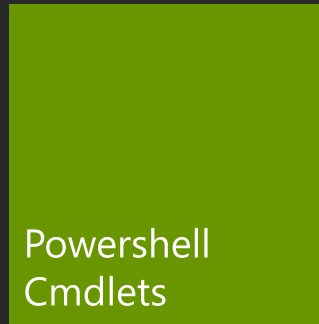
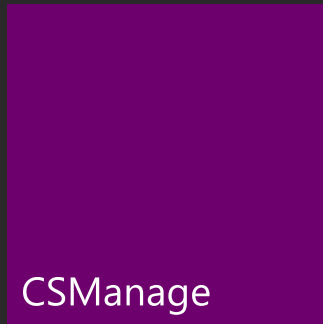
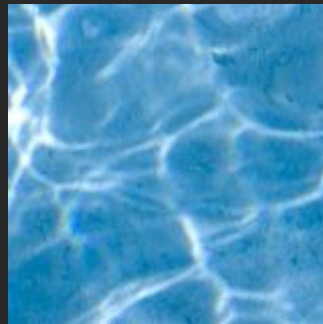
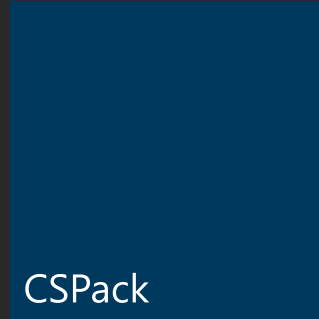
VIP swap



in place upgrade



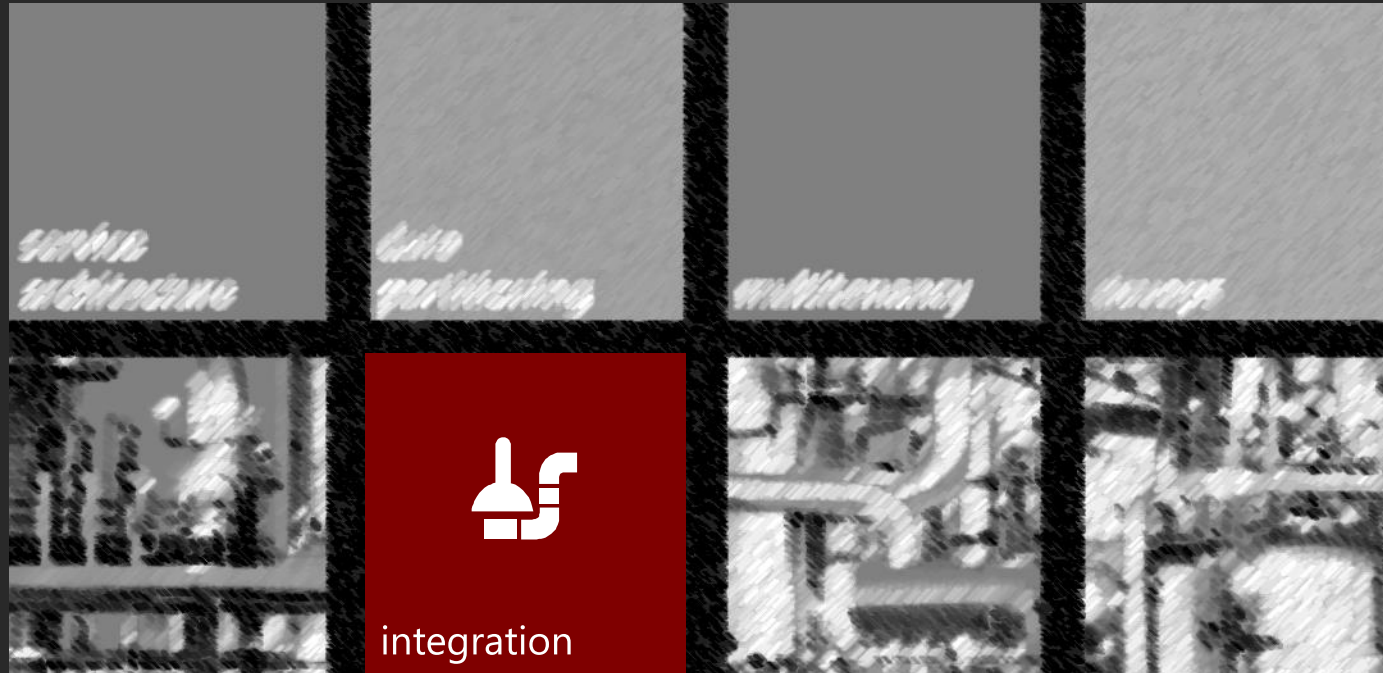
Windows Azure deployment tools




support offerings

	Entry-level reactive	Foundational reactive	Premium reactive	Premium reactive/proactive
Free	Web \$39/Month 6-month contract	Basic \$300/Month 6-month contract	Pro-Direct \$1K/Month 6-month contract	Premier \$3K/\$12K/\$44K/Month 1-yr contract
				Developer Mentoring & Proactive Services
			App Dev Mgr (Pooled) Azure Specific	App Dev Mgr (Designated) All Microsoft Products
			Unique Phone Line Elevated Priority	Unique Phone Line Elevated Priority MS Exec Escalations
		3 Phone Incidents per month	Unlimited Phone Incidents	Unlimited Phone Incidents
	Response Time Sev B: 4 Hrs Sev C: 8 hrs	Response Time Sev A: 2 hrs, Sev B: 4 Hrs Sev C: 8 hrs	Response Time Sev A: 1 hrs, Sev B: 2 Hrs Sev C: 4 hrs	Response Time Sev A: 1 hrs, Sev B: 2 Hrs Sev C: 4 hrs
	Unlimited Break/ Fix Support via Web	Unlimited Break/ Fix Support via Web	Unlimited Break/ Fix Support via Web	Unlimited Break/ Fix Support via Web

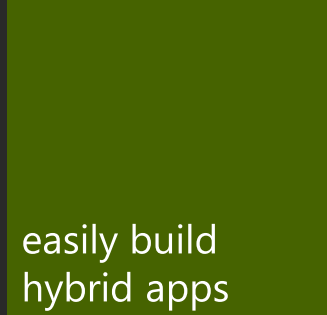
Community Forum | Service Dashboard and Outage Report | Billing and Account Support



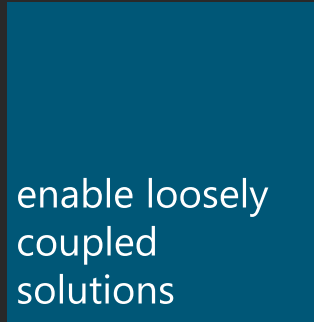
Windows Azure service bus



secure
messaging and
relay capabilities



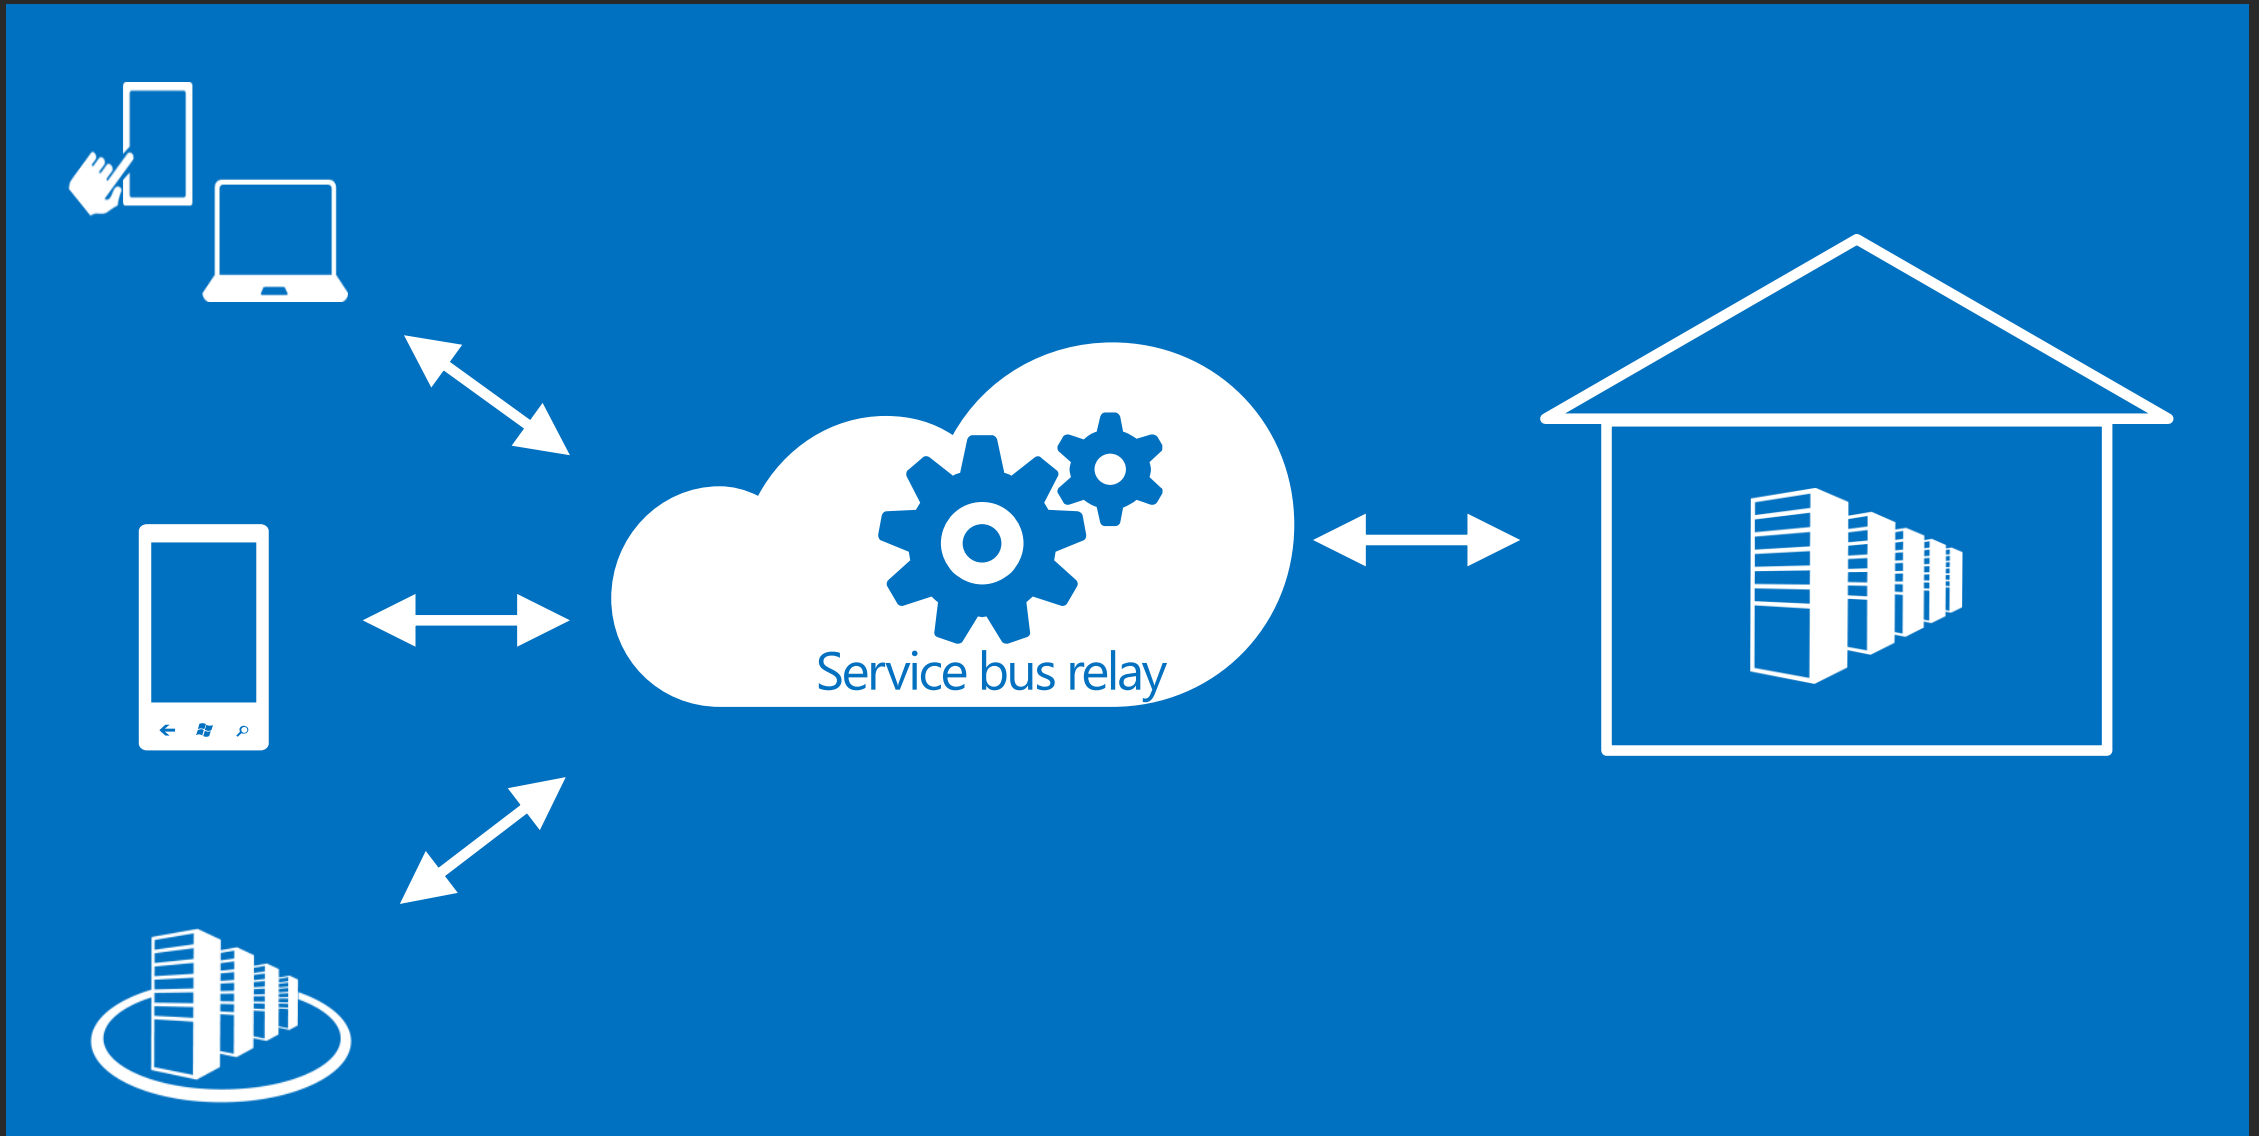
easily build
hybrid apps



enable loosely
coupled
solutions



accessing on-premise services





Demo:
Service Bus
Relay

conclusion



cloud style computing designs for scale out

scale out requires partitioning



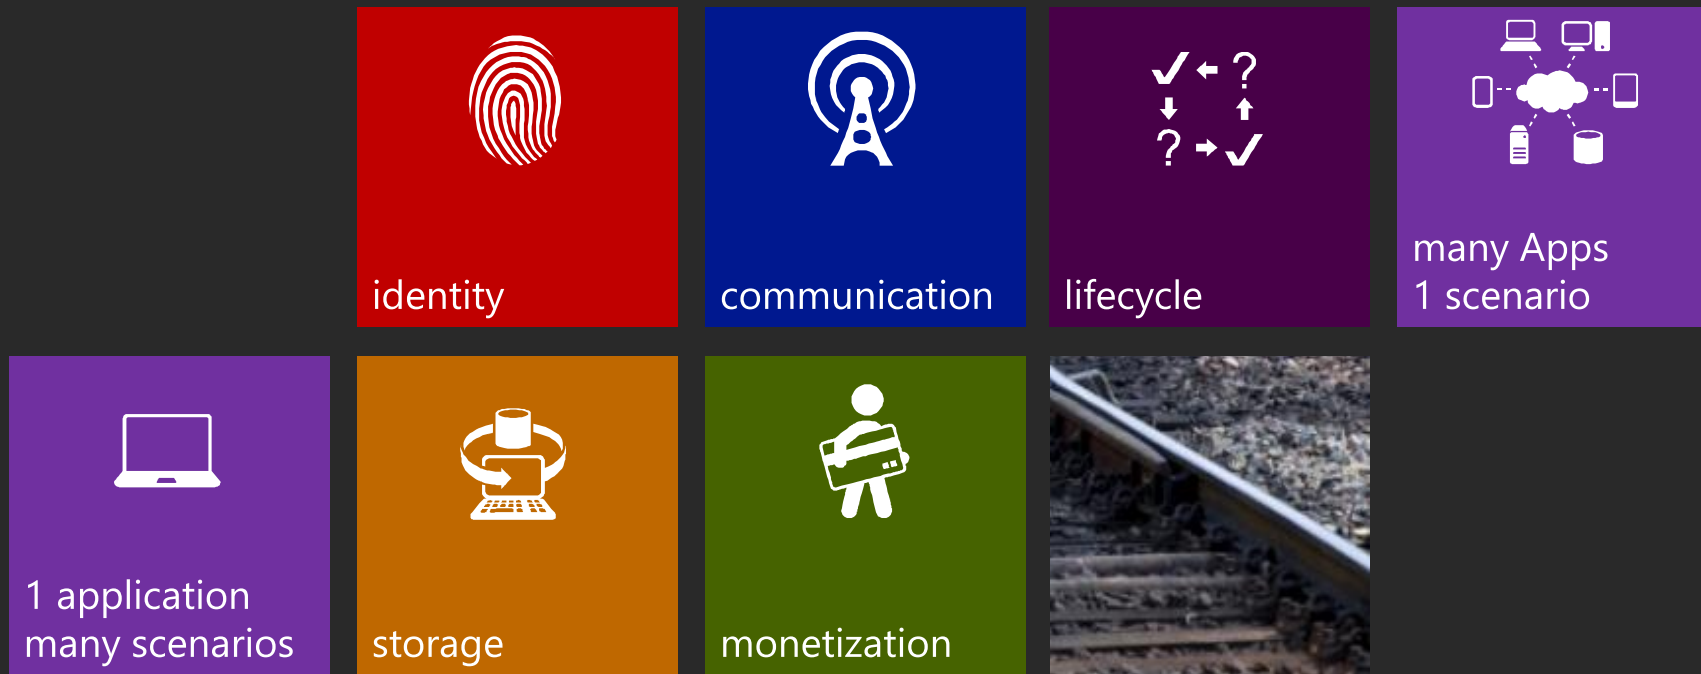
having redundant/duplicated data is ok

continuous delivery is a key asset



chapter V

from applications to apps



conclusion



from applications towards Apps

scenarios and tasks span multiple devices



cloud is a key enabler for connected devices

Windows Azure supports all devices

