cloud strategy day

"How the Microsoft cloud enables the best App experiences across devices"

Beat Schwegler – beatsch@microsoft.com
Director, Platform Strategy Group, Microsoft Corp.

agenda

Microsoft's devices and services strategy

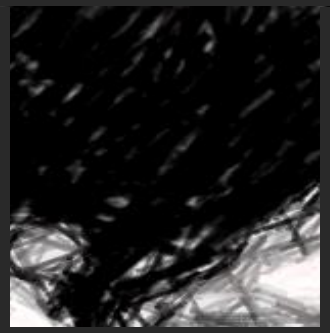

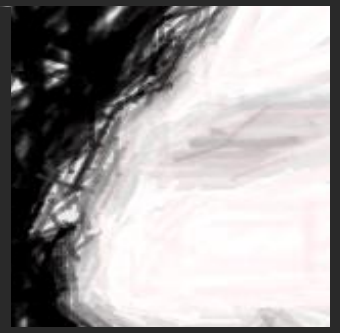

making money with apps

cloud enabled app scenarios with Windows Azure

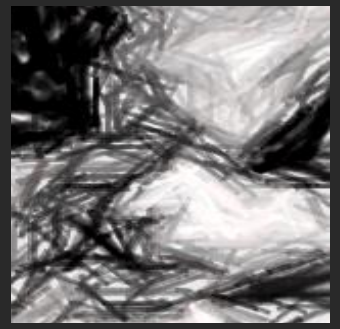

key architectural considerations

chapter IV


service
architecture


data
partitioning


multitenancy


DevOps


integration


service
architecture

service architecture

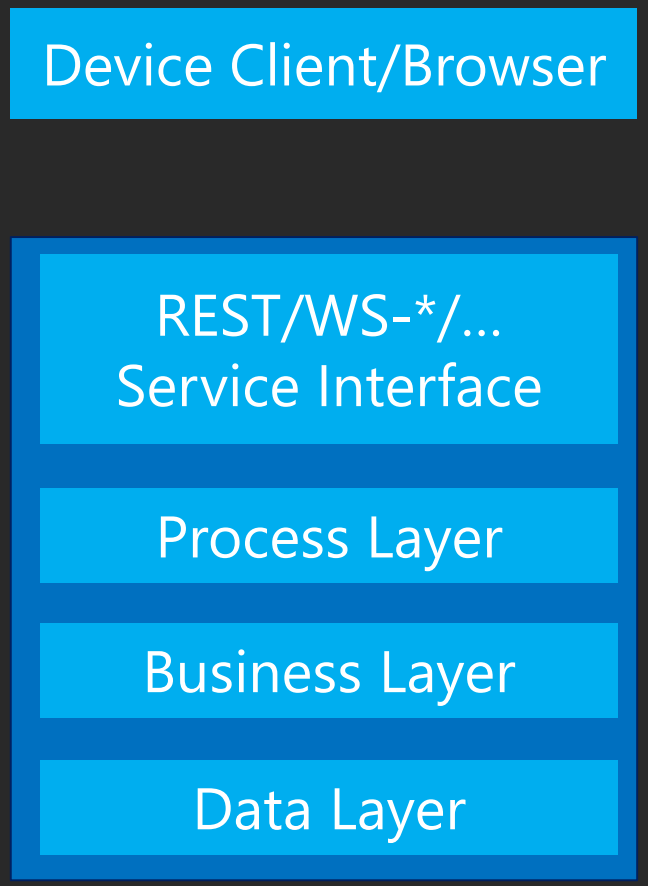

service architecture

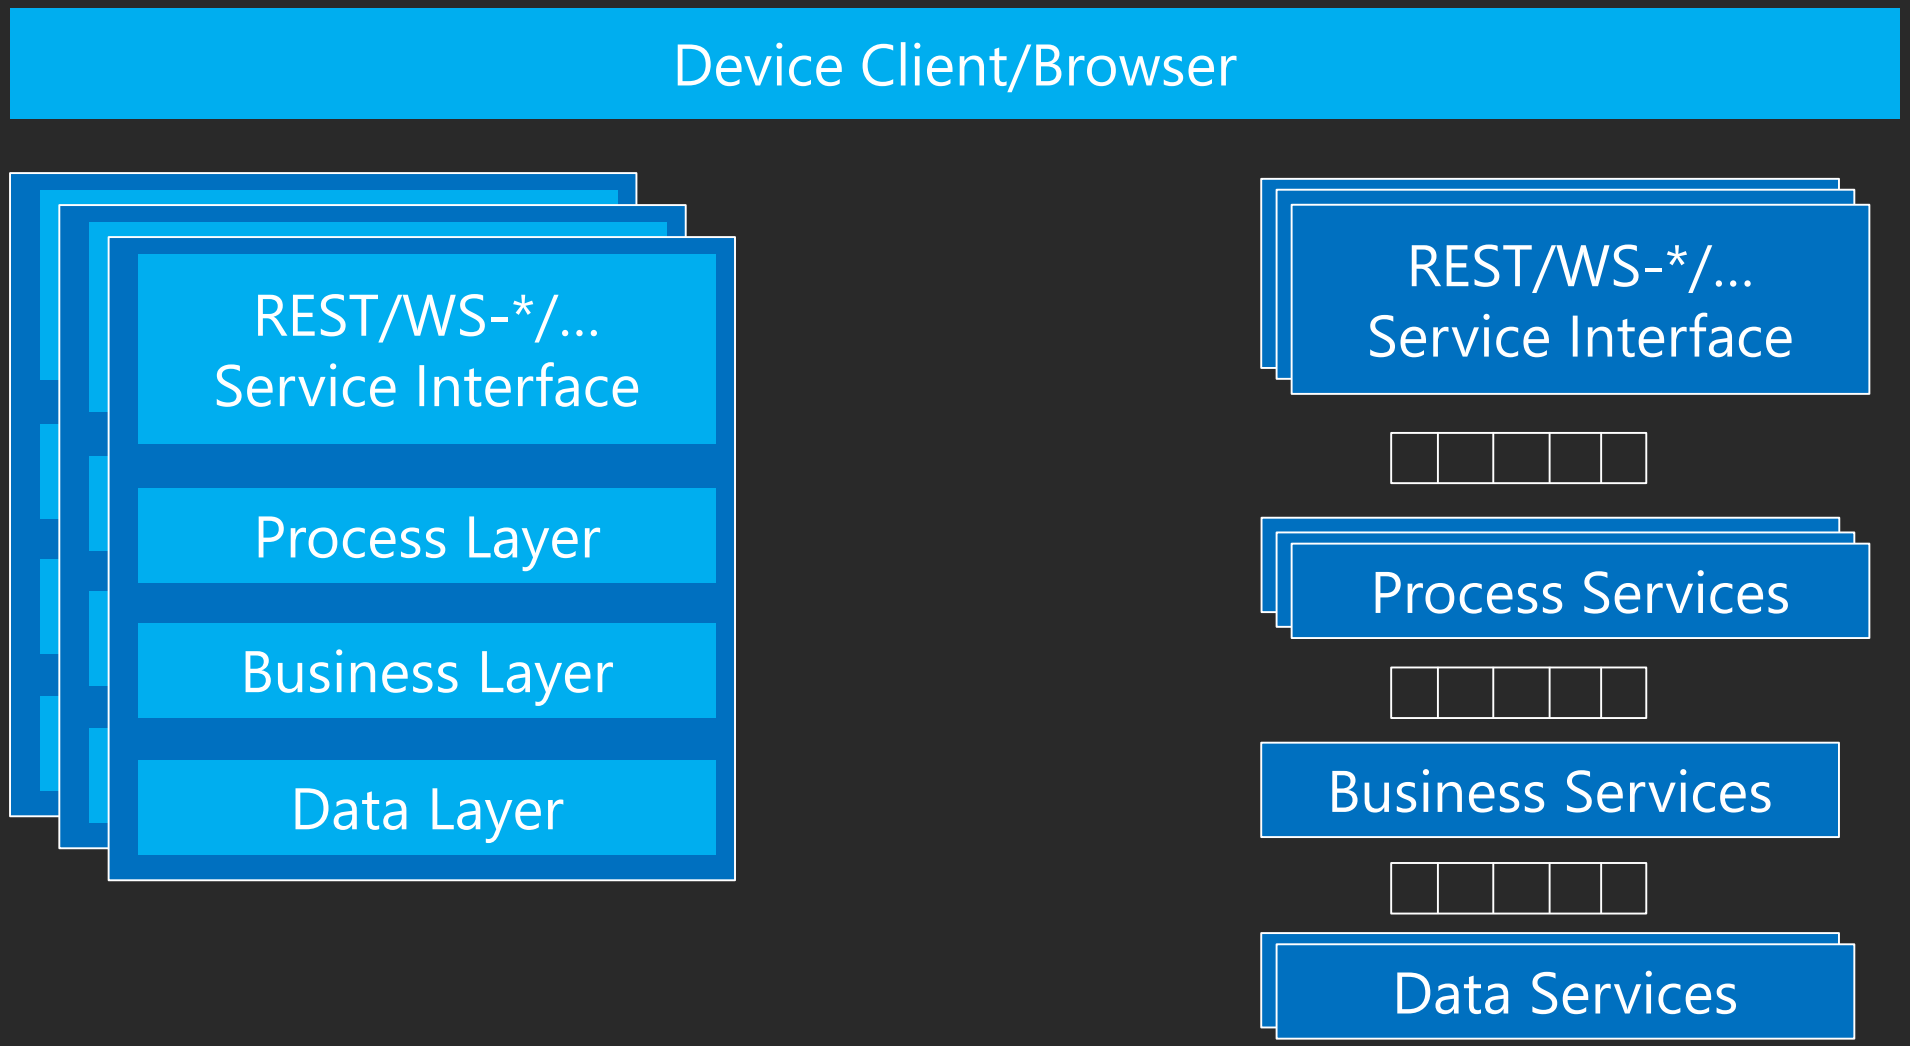

each instance contains all layers
inproc communication

instance per layer, communication through queues

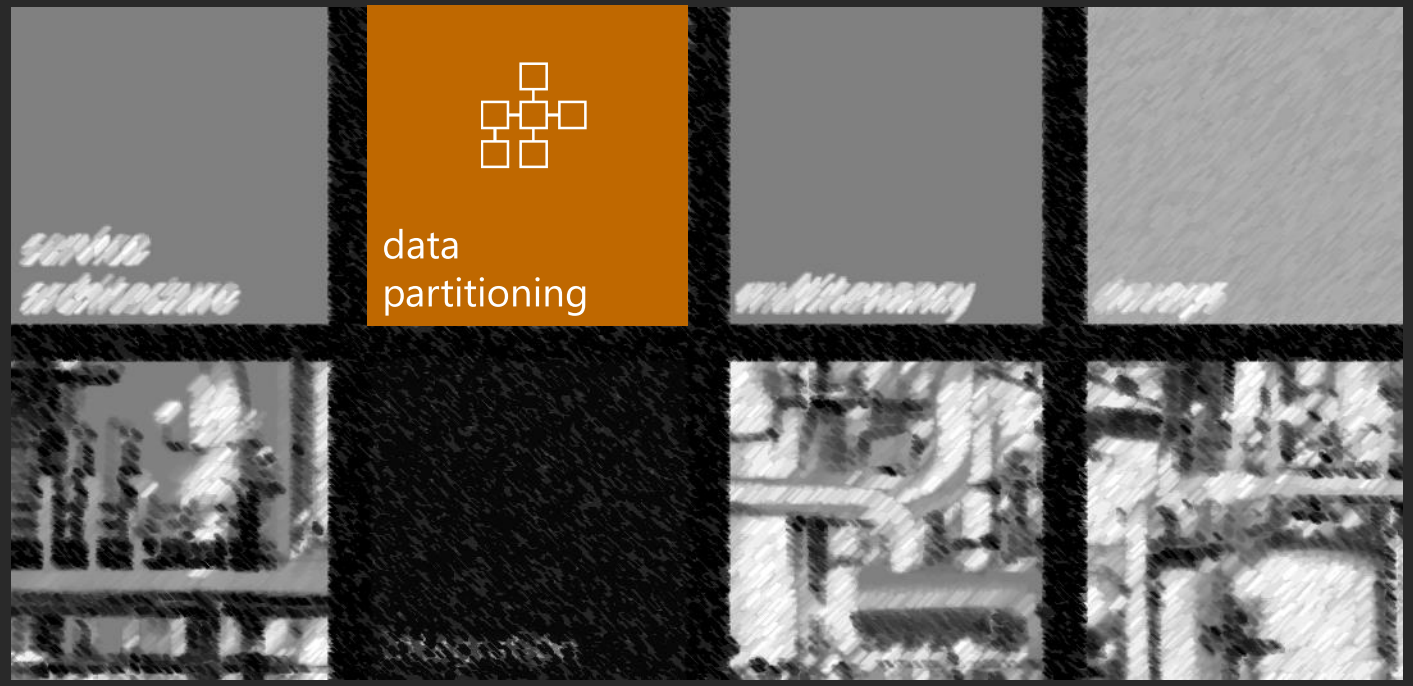
data
partitioning

cap theorem

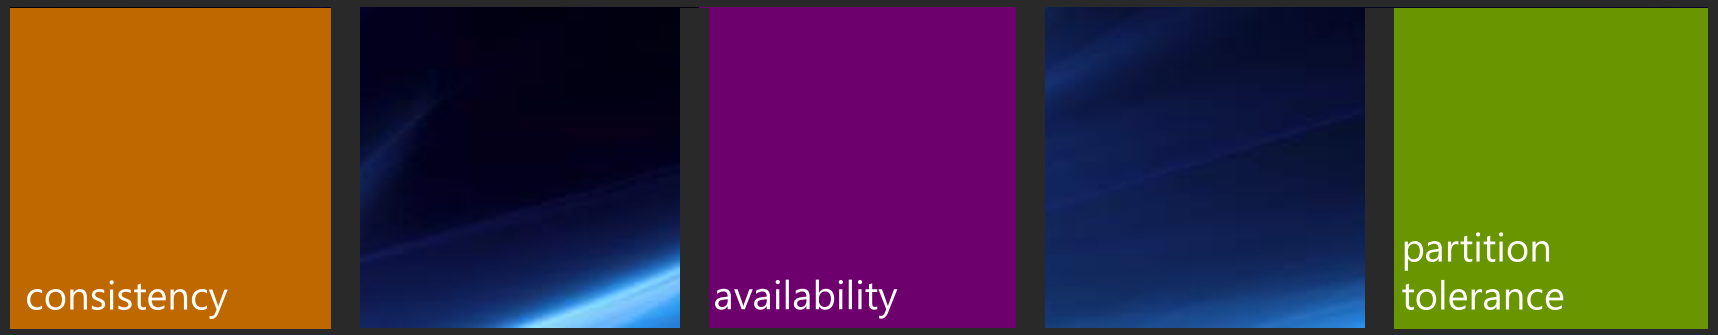

data partitioning

traditional reasons

data volume (too many bytes)

work load (too many transactions/second)

new 'cloud focused' reasons

cost (using different cost storage)

elasticity (just in time partitioning for high load periods)

horizontal partitioning

First Name	Last Name	Email	Thumbnail	Photo
David	Alexander	davida@contoso.com	3kb	3MB
Jared	Carlson	jaredc@contoso.com	3kb	3MB
Sue	Charles	suec@contoso.com	3kb	3MB
Simon	Mitchel	simonm@contoso.com	3kb	3MB
Richard	Zeng	richardz@contoso.com	3kb	3MB

horizontal partitioning (sharding)

spread data across similar nodes

achieve massive scale out (data and load)

intra-partition queries are simple

cross-partition queries are harder

vertical partitioning

First Name	Last Name	Email	Thumbnail	Photo
David	Alexander	davida@contoso.com	3kb	3MB
Jared	Carlson	jaredc@contoso.com	3kb	3MB
Sue	Charles	suec@contoso.com	3kb	3MB
Simon	Mitchel	simonm@contoso.com	3kb	3MB
Richard	Zeng	richardz@contoso.com	3kb	3MB

SQL Azure

Tables

BLOBS

vertical partitioning

place frequently queried data in more 'expensive' indexed storage

retrieving a whole row requires >1 query

spread data across dis-similar nodes

place large data in 'cheap' binary storage

hybrid partitioning

First Name	Last Name	Email	Thumbnail	Photo
David	Alexander	davida@contoso.com	3kb	3MB
Jared	Carlson	jaredc@contoso.com	3kb	3MB
Sue	Charles	suec@contoso.com	3kb	3MB
Simon	Mitchel	simonm@contoso.com	3kb	3MB
Richard	Zeng	richardz@contoso.com	3kb	3MB

tables != rdbms

cross partition queries are resource intensive

aggressive data duplication can save money and boost performance

storage is cheap

goal: To be able to include Partition Key in all queries

e.g. tweet storage

Tweet
TweetID
UserID
DateTimeStamp
Message

With an RDBMS you'd probably start something like this:
SELECT * FROM Tweet WHERE Message Like %SearchTerm%

e.g. tweet storage

You'd soon realize that LIKE isn't so wonderful.
You'd do a little normalization

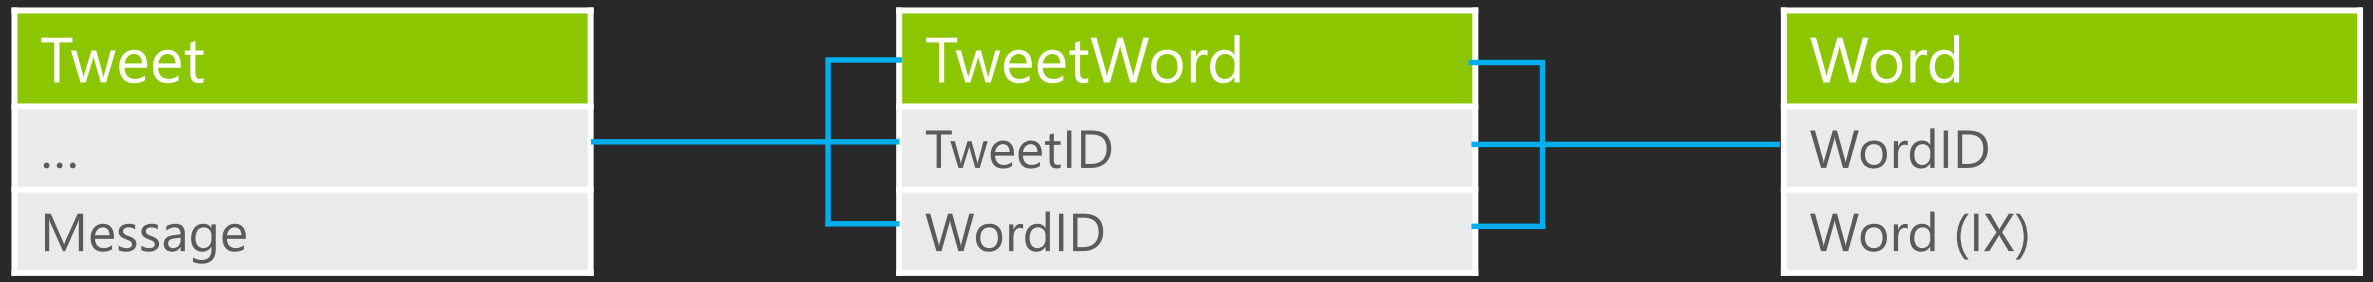

Which quickly becomes this as len(key) approaches AVG len(word)

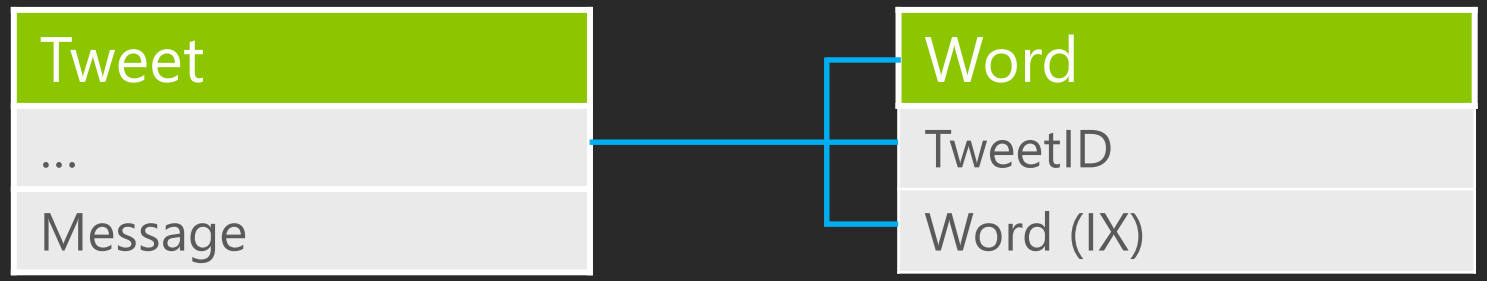

e.g. tweet storage

With Tables we go the whole way

Tweet
TweetID (RK)
UserID (PK)
DateTimeStamp
Message

Worker Role Creates

TweetIndex
TweetID (RK)
UserID
DateTimeStamp
Message
Word (PK)

GET All Entities in Partition ‘DavidA’ from Tweet

GET All Entities in Partition ‘Foo’ from TweetIndex

e.g. tweet storage

We may create multiple indexes

Tweet
TweetID (RK)
UserID (PK)
DateTimeStamp
Message

Worker Role Creates

MentionIndex
TweetID (RK)
UserID
DateTimeStamp
Message
MentionedUserID (PK)

GET All Entities in Partition ‘DavidA’ from MentionIndex

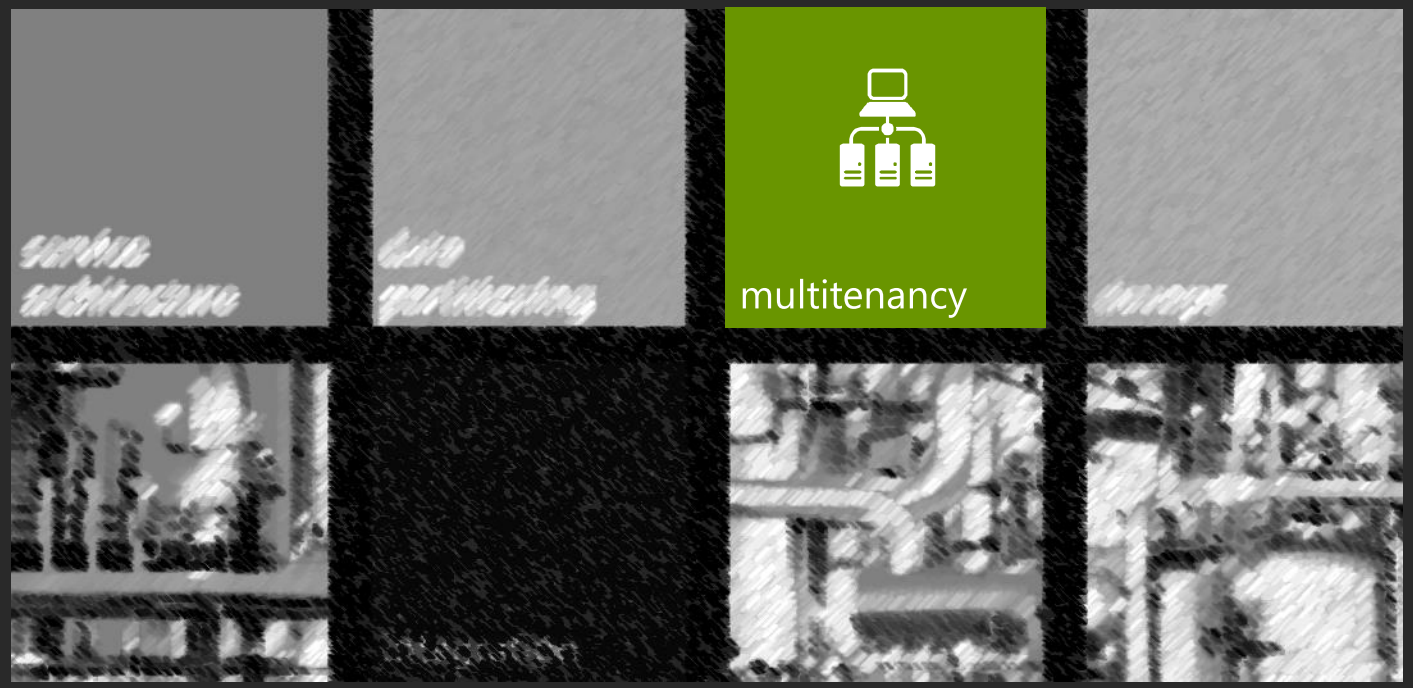
multitenancy

multitenancy

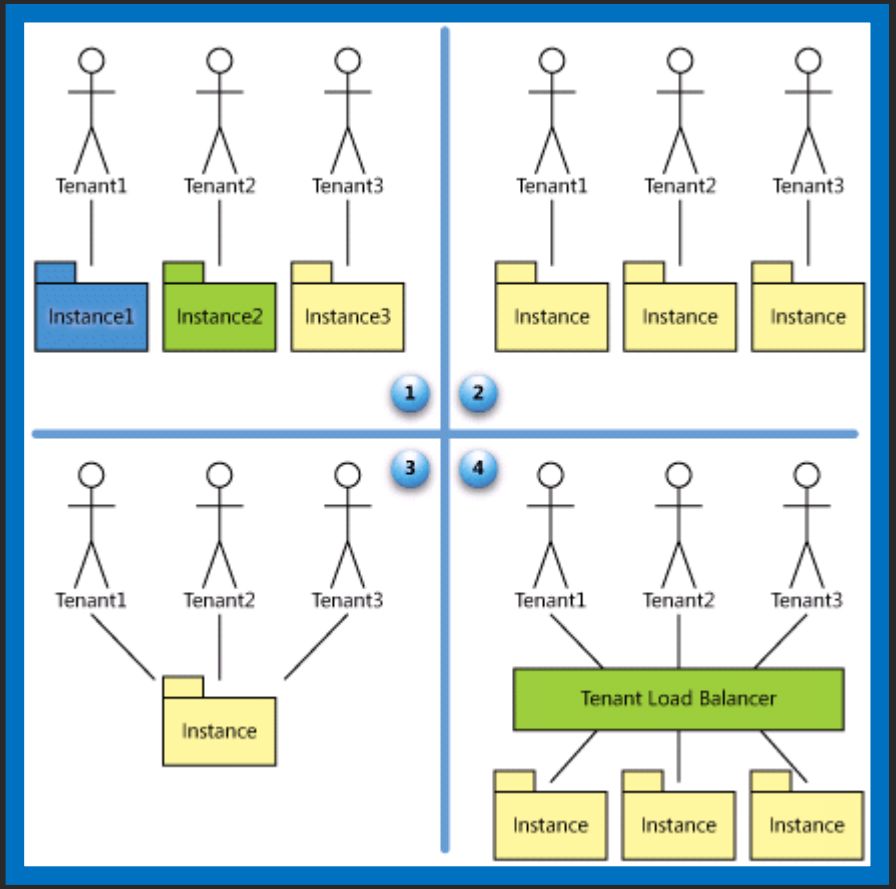


DevOps

continues deployment

http://code.flickr.com/

devops

development

quality
assurance

deployment

operations

one
team

continuous
delivery

devops

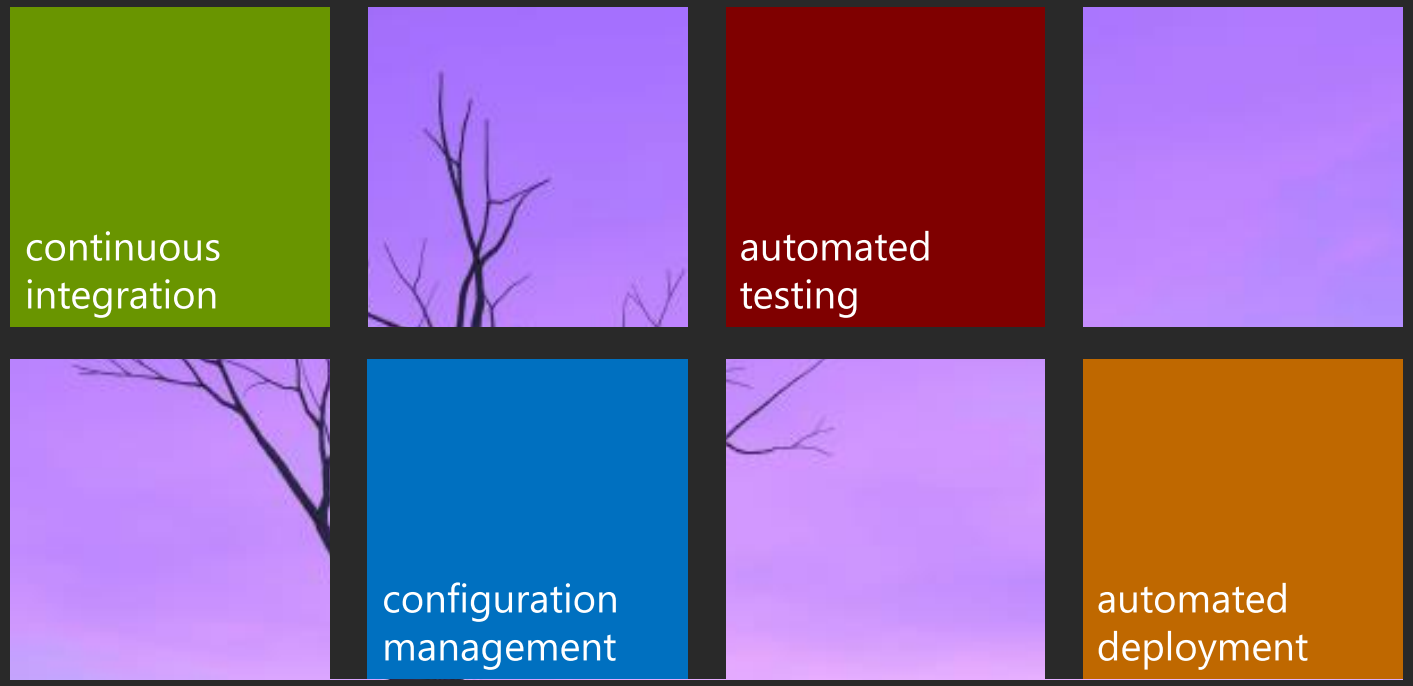
continuous integration
automated testing
configuration management
automated deployment

deployment/release strategies

upgrades in Windows Azure

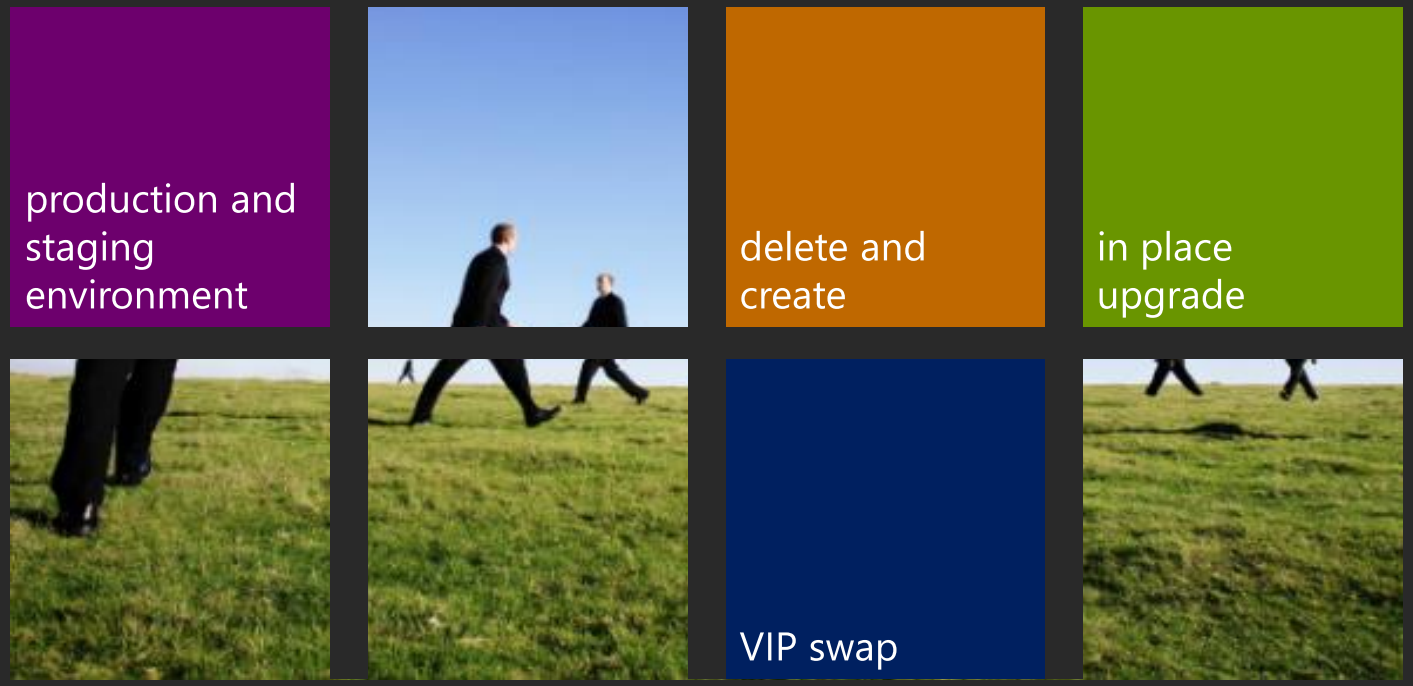

fault and upgrade domains

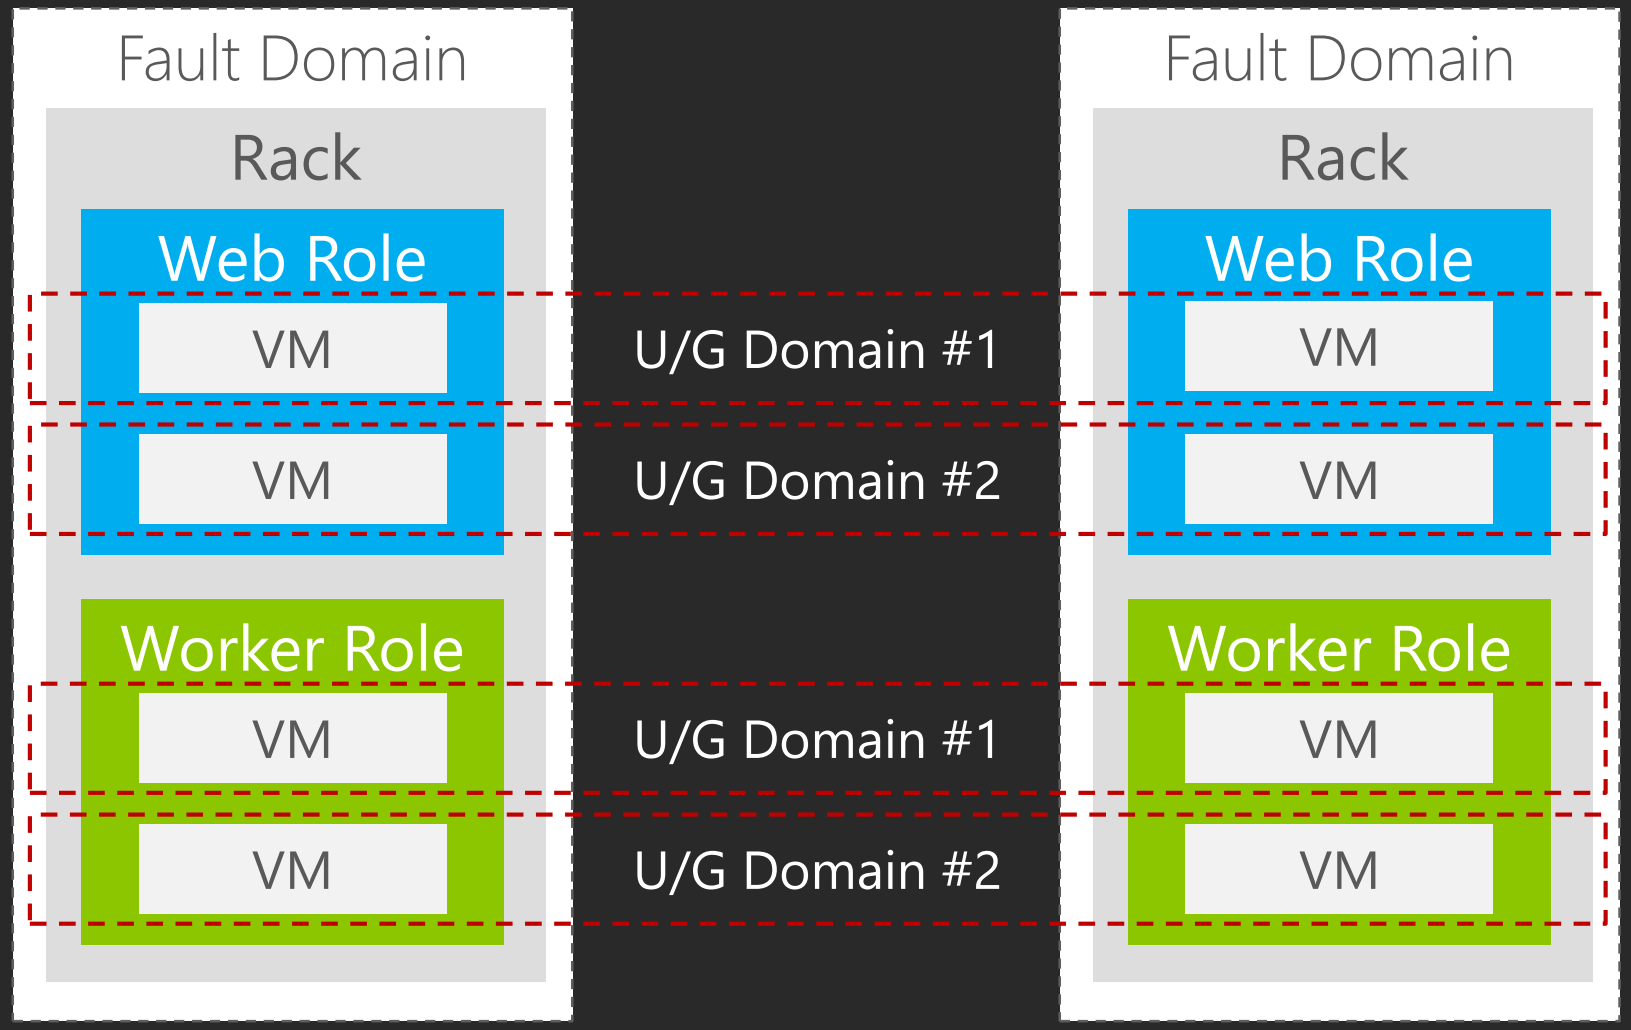

VIP swap

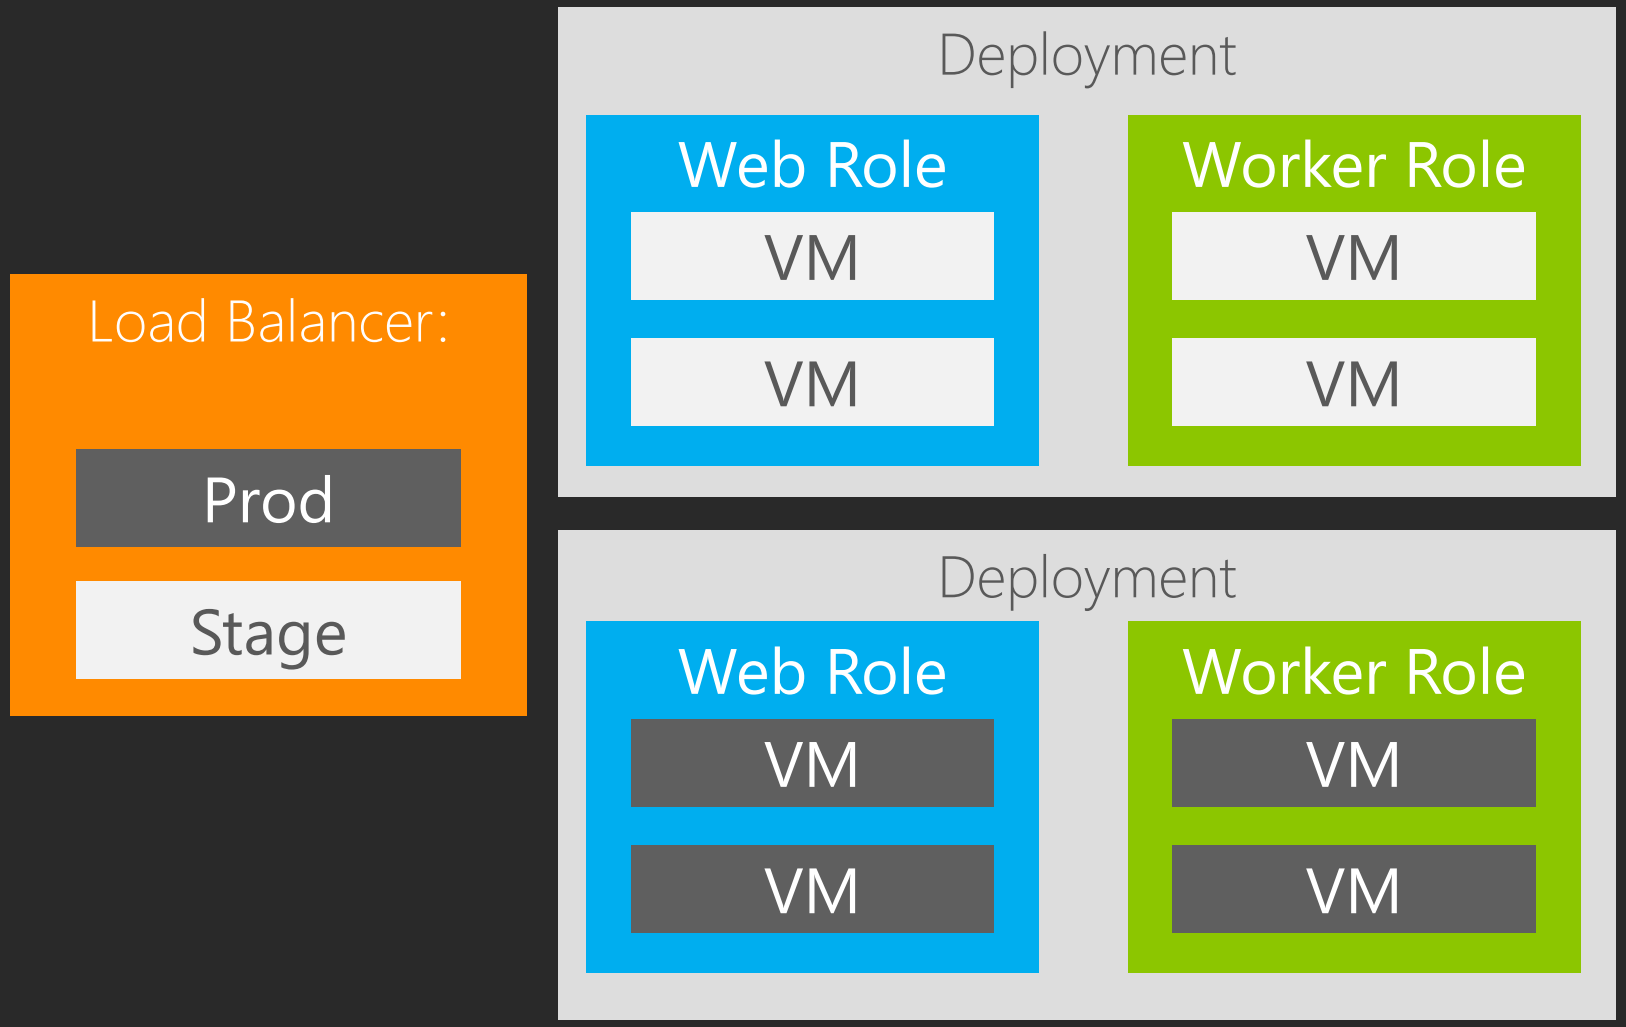
Load Balancer:
Prod
Stage
Deployment
Web Role
VM
VM
Worker Role
VM
VM
Deployment
Web Role
VM
VM
Worker Role
VM
VM

in place upgrade

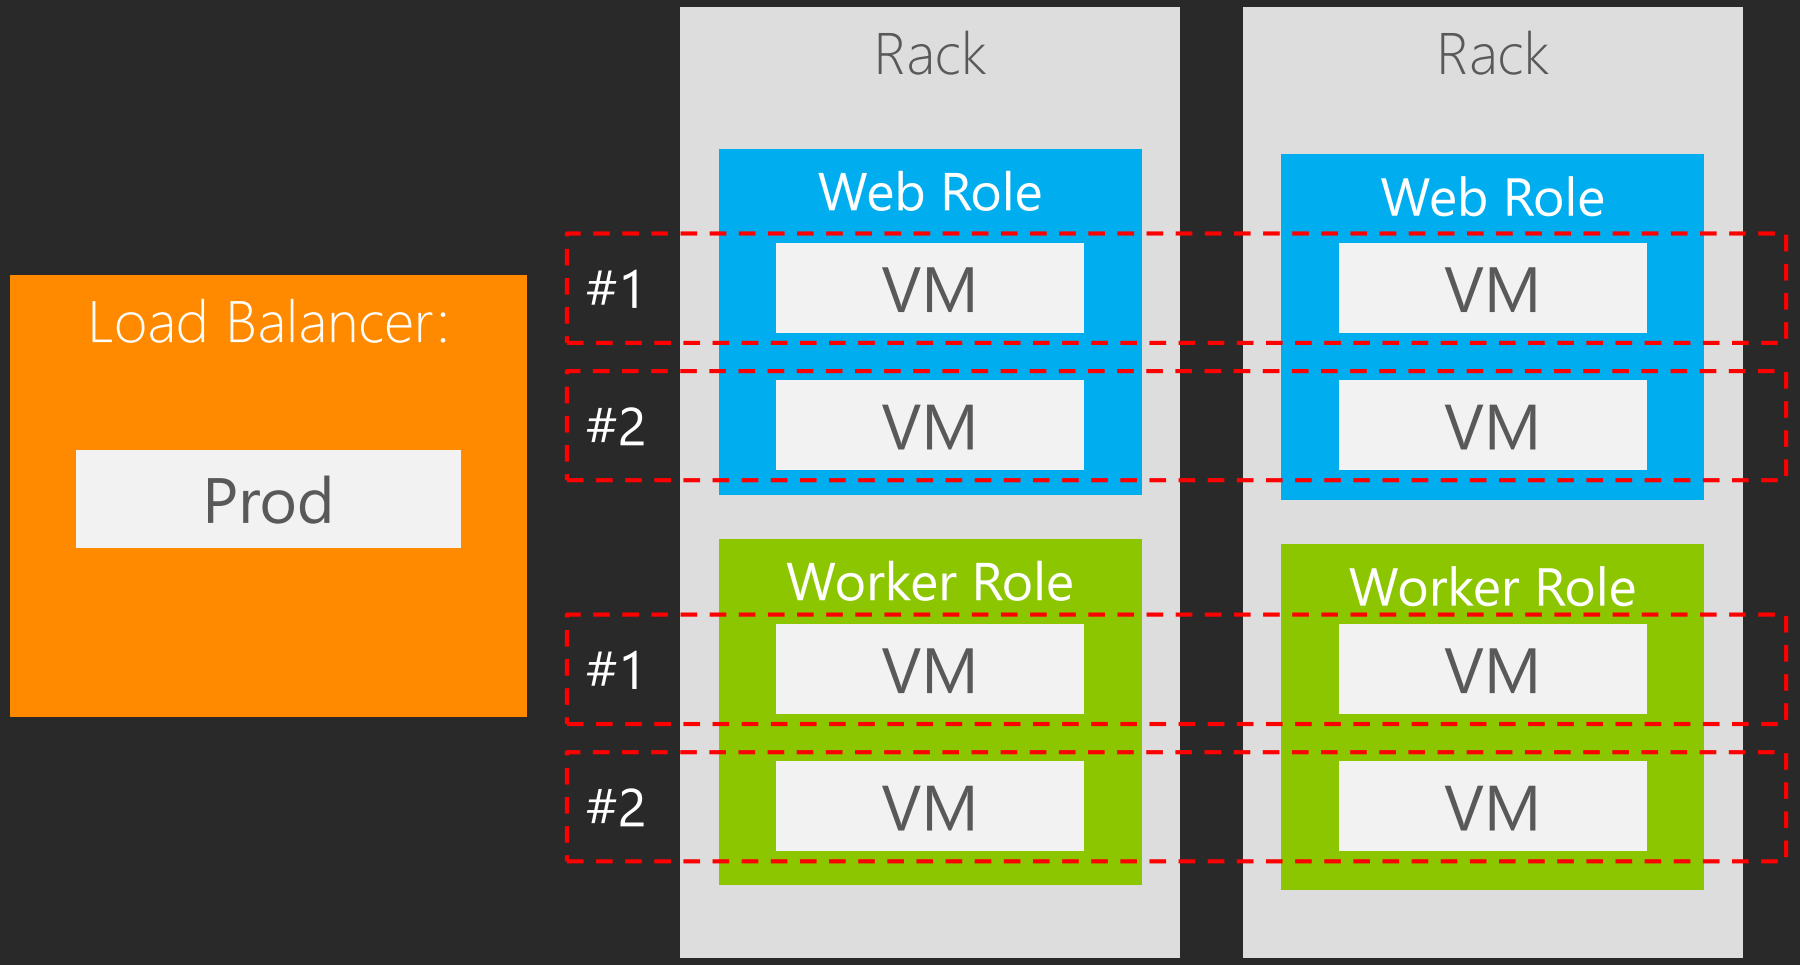

Windows Azure deployment tools

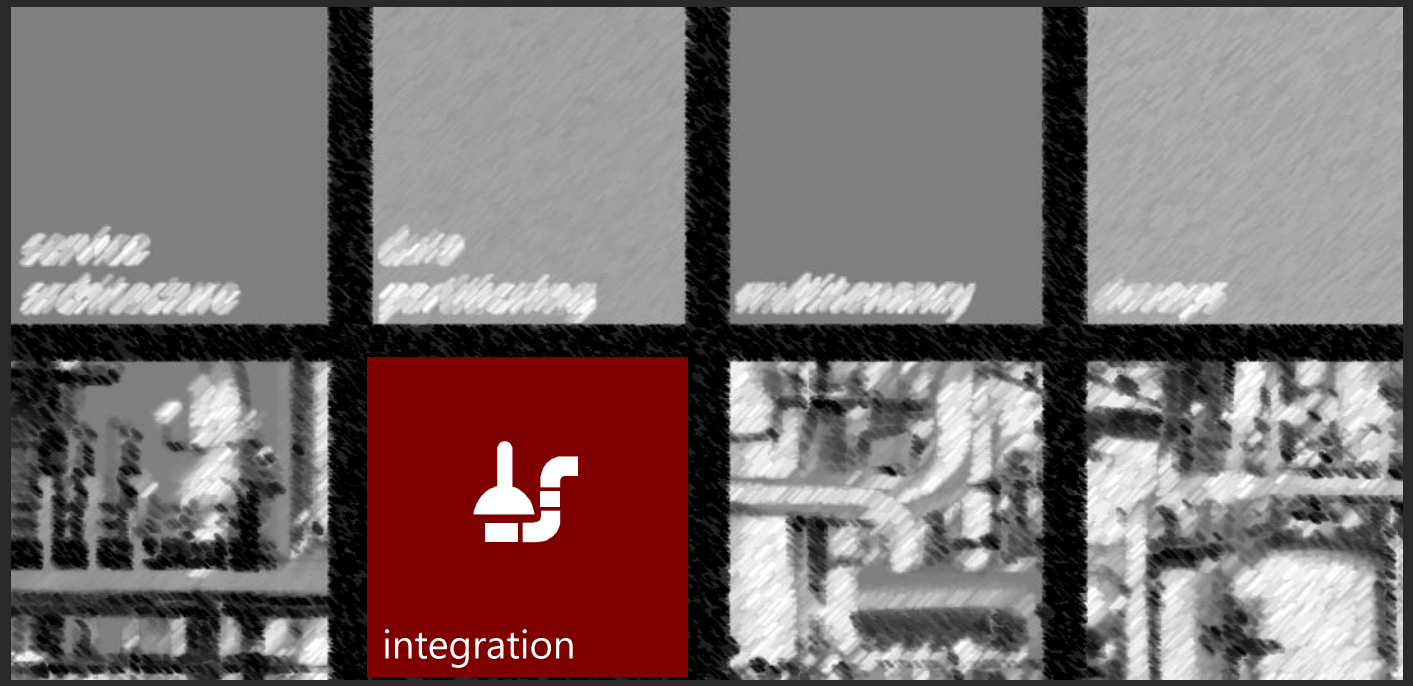
integration

Windows Azure service bus

secure messaging and relay capabilities

easily build hybrid apps

enable loosely coupled solutions

accessing on-premise services

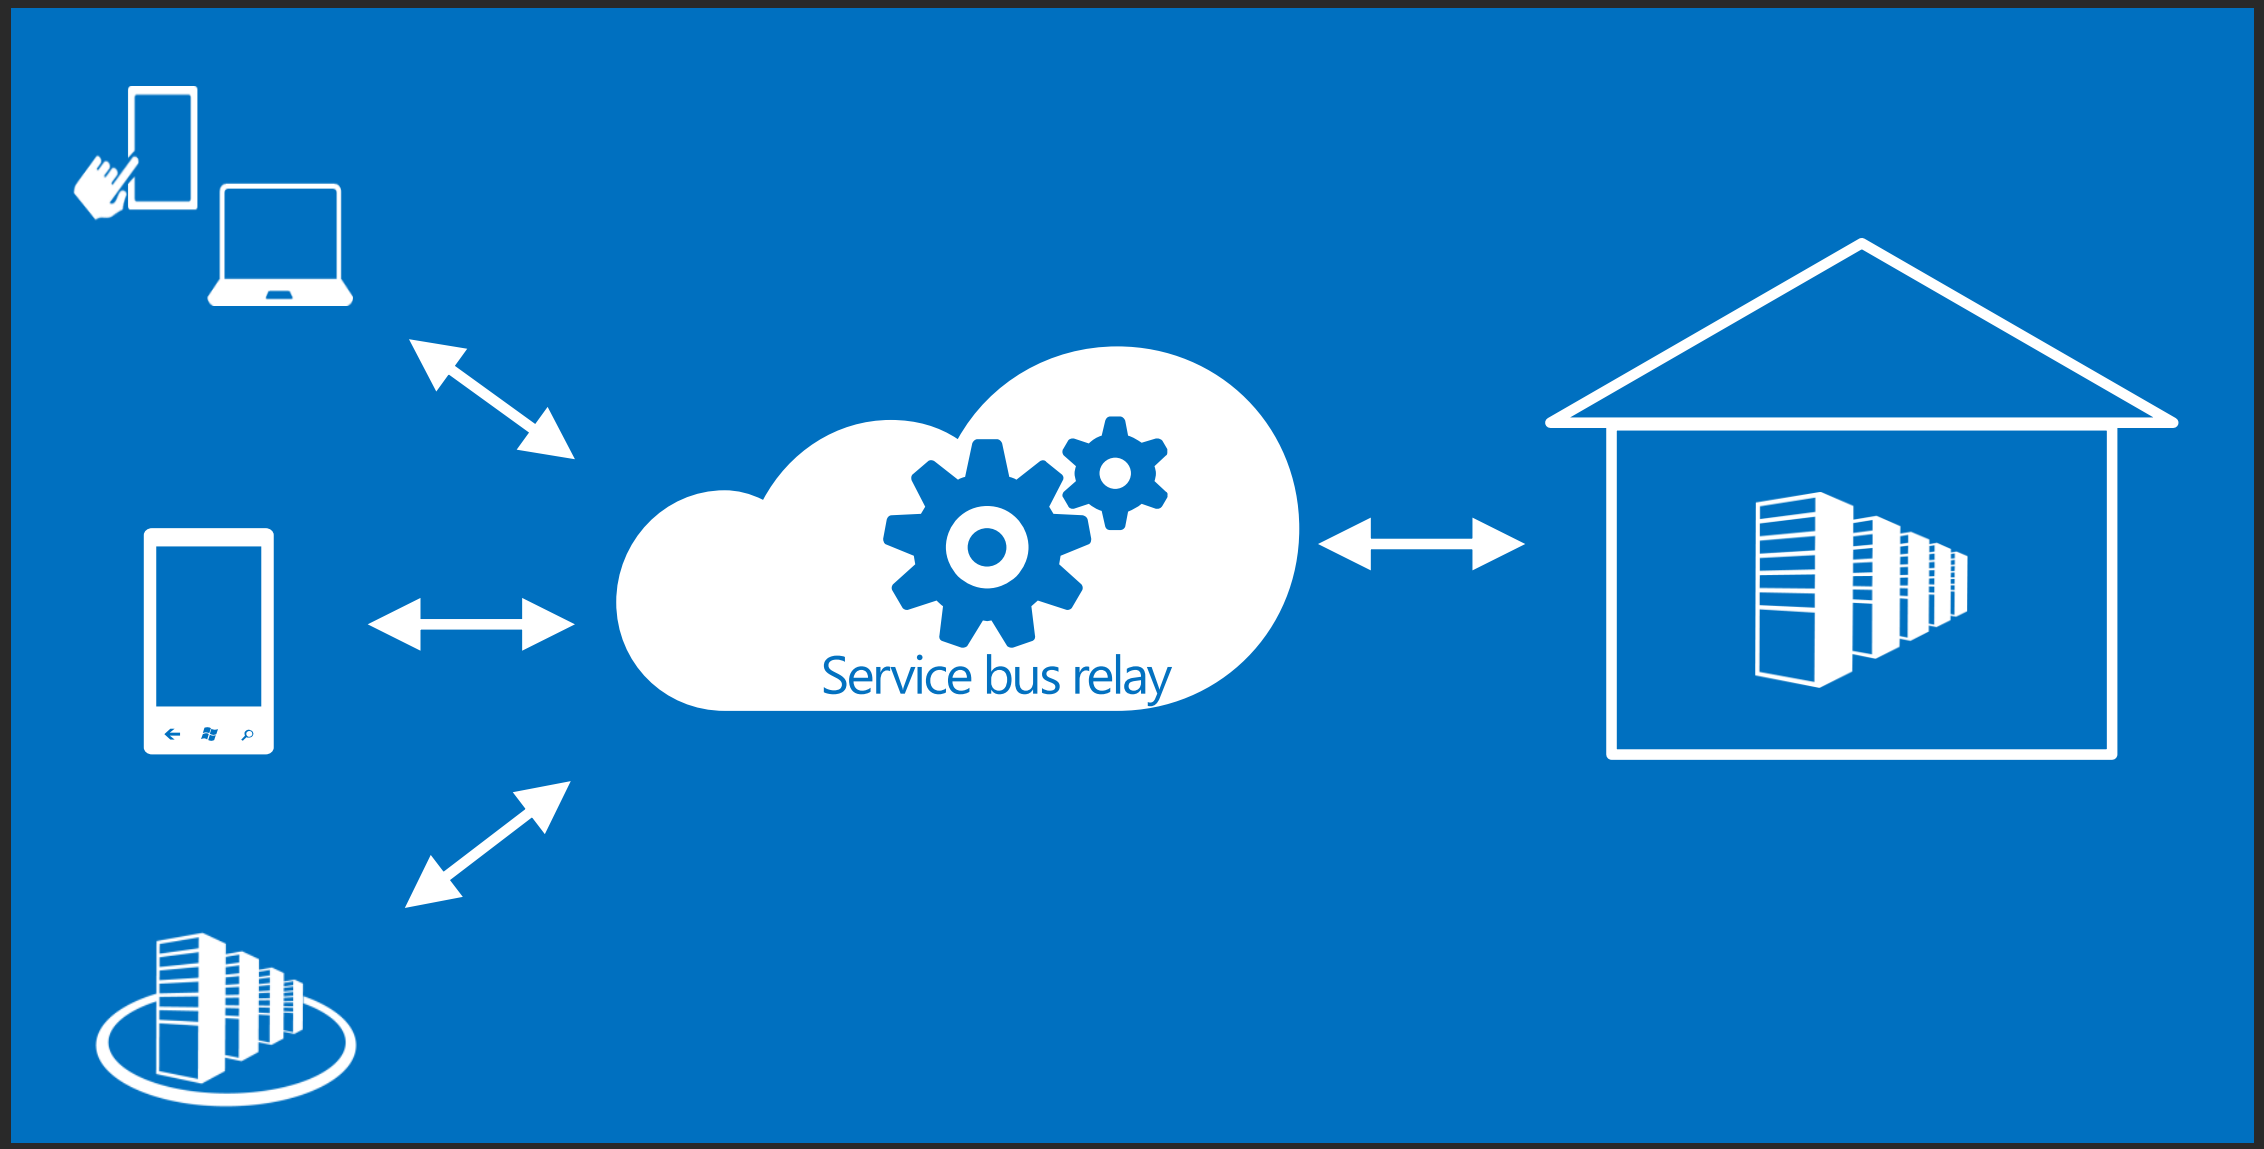

Demo:
Service Bus
Relay

conclusion

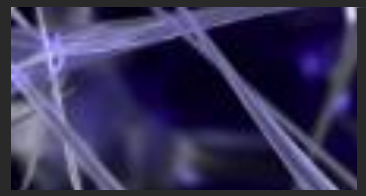

cloud style computing designs for scale out

scale out requires partitioning

having redundant/duplicated data is ok

continuous delivery is a key asset

chapter V

from applications to apps

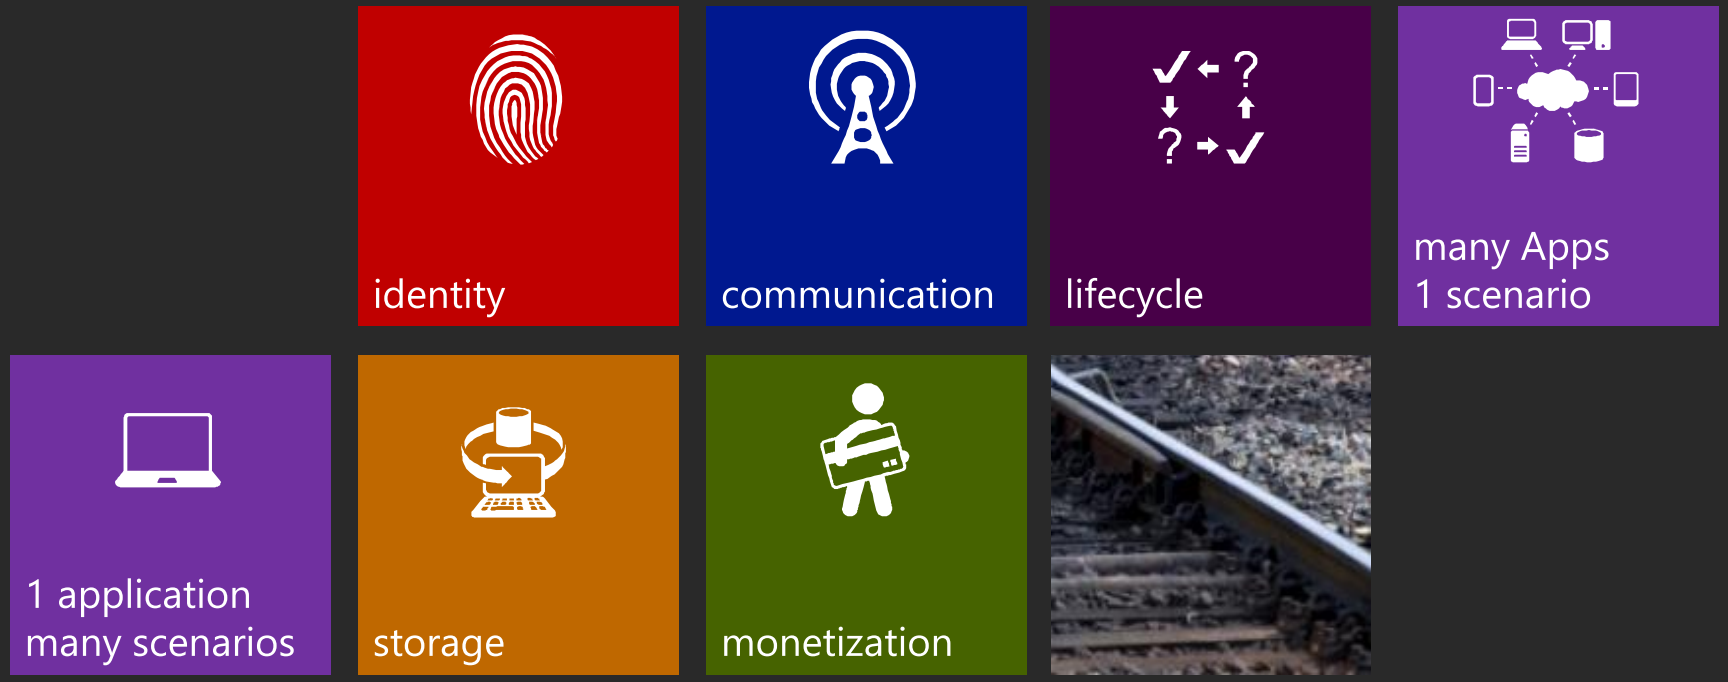

conclusion

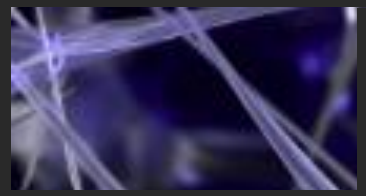

from applications towards Apps

scenarios and tasks span multiple devices

cloud is a key enabler for connected devices

Windows Azure supports all devices